

# Working with Data

Fraida Fund

## Contents

|                                                                          |    |
|--------------------------------------------------------------------------|----|
| Recap: Machine learning is data driven . . . . .                         | 2  |
| Garbage in, garbage out . . . . .                                        | 2  |
| Model training vs evaluation vs deployment . . . . .                     | 3  |
| Working with data: two stages . . . . .                                  | 4  |
| Stage 1: Is this data appropriate for the task? . . . . .                | 4  |
| Legal and ethical concerns . . . . .                                     | 4  |
| Representativeness concerns . . . . .                                    | 5  |
| Predictive feature concerns . . . . .                                    | 6  |
| Example: distinguish real headshots from AI generated headshots. . . . . | 7  |
| Other data quality concerns . . . . .                                    | 8  |
| Stage 2: How do I process the data? . . . . .                            | 8  |
| Make and check assumptions . . . . .                                     | 8  |
| Split data . . . . .                                                     | 9  |
| Select features . . . . .                                                | 11 |
| Select target . . . . .                                                  | 11 |
| Construct samples . . . . .                                              | 12 |
| Handle missing data . . . . .                                            | 12 |
| Convert to numeric types . . . . .                                       | 13 |
| Create “transformed” features . . . . .                                  | 16 |
| Recap: Working with data . . . . .                                       | 16 |

## Recap: Machine learning is data driven

- Rules based: you can write rules to produce an output using only the current input.
- Data driven: to produce an output/write rules to produce an output, you need to use *other data samples* besides for the current input.

For example, we introduced *prediction by mean* as the simplest data-driven “model”.

If we select a model class where we make the same prediction,  $w$ , for every sample:

$$\hat{y} = w$$

and we use a mean squared error loss function, where we want to minimize:

$$\frac{1}{N} \sum (y - \hat{y})^2$$

we can show that the optimal value of the parameter  $w$ , is the mean of  $y$  across all samples.

Step 1: plug  $\hat{y} = w$  into the loss function:

$$\text{minimize } \frac{1}{N} \sum (y - \hat{y})^2 \rightarrow \text{minimize } \frac{1}{N} \sum (y - w)^2$$

Step 2: take the derivative with respect to the parameter  $w$

$$\frac{\partial}{\partial w} \frac{1}{N} \sum (y - w)^2 = -\frac{2}{N} \sum (y - w)$$

Step 3: set the derivative equal to 0, and solve for the parameter  $w$

$$-\frac{2}{N} \sum (y - w) = 0 \quad \text{Set the derivative equal to zero}$$

$$\sum (y) - \sum (w) = 0 \quad \text{Multiply by } -\frac{N}{2} \text{ and expand}$$

$$\sum (y) = \sum (w) = Nw \quad \text{Sum the constant } w \text{ over } N \text{ samples}$$

$$\rightarrow w = \frac{1}{N} \sum (y) \quad \text{Solve for } w$$

## Garbage in, garbage out

Any machine learning project has to start with “good” data.

There is a “garbage in, garbage out” rule: If you use “garbage” to train a machine learning model, you will only get “garbage” out.

And: Since you are evaluating on the same data, you might not even realize it is “garbage” at first! You may not realize until the model is already deployed in production! This is the absolute worst case scenario, as illustrated below...

## Model training vs evaluation vs deployment




Figure 1: The lifecycle of an ML model

We want to understand how the model will behave in *deployment* as early as possible (before investing too much time, effort, money in a model that won't do well).

| Training Accuracy | Evaluation Accuracy | Deployment Accuracy | Outcome |
|-------------------|---------------------|---------------------|---------|
| 0.95              | 0.93                | 0.91                | 😄       |
| 0.55              | 0.52                | N/A                 | 😐       |
| 0.95              | 0.93                | 0.51                | 😱       |

- Best case: Model does well in evaluation, and deployment
- Second best case: Model does poorly in evaluation, is not deployed
- Worst case: Model does well in evaluation, poorly in deployment (“overly optimistic evaluation”)

When preparing to work with data, then, we have to think about both of the ways in which our data will be used:

- to *model* a relationship between  $\hat{y} = f(x)$ . This stage is designed to optimize the training objective, *on the training data*.
- and then to *evaluate* the model  $f(x)$ , to see how well  $\hat{y}$  agrees with  $y$ . We evaluate on a separate held-out test set not used in training. (It's typical for the performance to be slightly worse on the test set, since the model is optimized on the training set!)

The goal of the evaluation is to estimate how the model will perform in deployment, which is what we *really* care about and want to optimize.

We need to make sure that we have suitable data for *both* purposes.

## Working with data: two stages

1. Is this data appropriate for the task?
2. How should I prepare it for modeling?

### Stage 1: Is this data appropriate for the task?

#### Legal and ethical concerns

- **Consent:** Did people agree to this data being collected and used for this purpose?
- **Privacy:** Could the data identify people or reveal sensitive information?
- **Copyright:** Do we have the right to use, modify, and redistribute the data?
- **Bias/fairness:** Could it produce unequal outcomes?

Before we ask whether a dataset is useful, we need to ask whether we legally may or ethically should use it at all. Otherwise we risk expensive lawsuits, or being shut down by regulators.

Here is one example of a failure related to each concern:

- **Consent:** Everalbum used customers' photos to develop facial recognition technology without obtaining permission. In 2021, the FTC required the company to delete the photos and facial recognition models.
- **Privacy:** OkCupid operated a dating platform whose users shared sensitive information, including sexual orientation, location, and personal responses. An individual scraped about 70,000 profiles and published the data without users' consent, making people vulnerable to identification even if their names were removed. In 2016, the disclosure triggered public backlash and forced the dataset offline.
- **Copyright:** Stability AI used millions of Getty images to train Stable Diffusion without permission, according to Getty Images. In 2023, Getty sued Stability AI, which faced costly copyright litigation in the United States and United Kingdom.
- **Bias/fairness:** Facebook used a machine learning ad delivery system to optimize who received housing ads, and the system let advertisers exclude people based on protected characteristics. In 2019, HUD charged Facebook with violating the Fair Housing Act; Meta later agreed to restrict housing ad targeting and change its delivery system.

and some citations/further reading:

- Marin Benčević, Marija Habijan, Irena Galić, Danilo Babin, Aleksandra Pižurica, Understanding skin color bias in deep learning-based skin lesion segmentation, Computer Methods and Programs in Biomedicine, Volume 245, 2024, <https://doi.org/10.1016/j.cmpb.2024.108044>
- Ziad Obermeyer, Brian Powers, Christine Vogeli, Sendhil Mullainathan, Dissecting racial bias in an algorithm used to manage the health of populations, Science, Volume 366, Issue 6464, 2019, pp. 447–453, <https://doi.org/10.1126/science.aax2342>
- Shirin Ghaffary, The algorithms that detect hate speech online are biased against black people, Vox, August 15, 2019, <https://www.vox.com/recode/2019/8/15/20806384/social-media-hate-speech-bias-black-african-american-facebook-twitter>
- Benj Edwards, Artist finds private medical record photos in popular AI training data set, Ars Technica, September 21, 2022, <https://arstechnica.com/information-technology/2022/09/artist-finds-private-medical-record-photos-in-popular-ai-training-data-set/>
- Joseph Cox, "70,000 OkCupid Users Just Had Their Data Published," Vice, May 12, 2016. <https://www.vice.com/en/article/70000-okcupid-users-just-had-their-data-published/>
- Geoffrey A. Fowler, Millions of people uploaded photos to the Ever app. Then the company used them to develop facial recognition technology, NBC News, May 9, 2019. <https://www.nbcnews.com/tech/security/millions-people-uploaded-photos-ever-app-then-company-used-them-n1003371>
- David Meyer, Amazon reportedly killed an AI recruitment system because it couldn't stop the tool from discriminating against women, Fortune via Yahoo Finance, October 10, 2018. <https://finance.yahoo.com/news/amazon-reportedly-killed-ai-recruitment-100042269.html>

## Representativeness concerns

Is the data similar to deployment setting in:

- **Population:** Do the people, places, or objects in the data match the population where the model will be used?
- **Context/setting:** Does the location and data collection environment match where the model will be used?
- **Time period:** Does the data reflect the period when the model will make predictions?
- **Label balance:** Is the distribution of the target similar to what the model will encounter?

Representativeness affects both **modeling** and **evaluation**:

- During modeling, the model learns patterns from the population, setting, time period, and class proportions in the training data.
- During evaluation, those same factors determine whether the test score predicts performance after deployment. A model can receive a high score on a test set that does not resemble the real deployment population.

Here are some examples from practice:

- **Population:** The **Framingham cardiovascular risk score** was designed to estimate a patient's ten-year risk of coronary heart disease from factors such as age, sex, blood pressure, cholesterol, and smoking. Researchers trained it using data from the mostly white residents of Framingham, Massachusetts. When clinicians applied the score to other ethnic groups, its accuracy varied because those populations had different risk patterns. A model can therefore learn relationships that work for the training population but misestimate risk for people outside it.
- **Context/setting:** **Google Health's deep learning system** for detecting diabetic retinopathy from retinal photographs was deployed in clinics in Thailand. The system performed well under controlled evaluation conditions, but clinics used different lighting, camera conditions, and patient preparation practices, such as whether patients' pupils were dilated. In the first six months, 21% of 1,838 images were judged ungradable because they failed the model's quality requirements. The result showed how changes between development and deployment environments can reduce the practical reliability of an image classification system.
- **Time period:** **Google Flu Trends** was designed to estimate flu prevalence faster than official health reports by using patterns in Google searches. Google trained and calibrated the system against historical CDC flu data and search queries. After deployment, it overestimated flu cases, especially during the 2012-2013 flu season, because media coverage changed search behavior and Google changed its search system. The search patterns changed over time, so the model's original relationship between searches and flu cases no longer held.
- **Label balance:** An AI writing detector is designed to classify text as human written or AI generated. **Turnitin's detector** was trained on examples of both types of writing, but the true proportion of AI generated text in student submissions was not known and would not necessarily match the training proportion. When a rare class appears much more often in training or evaluation than it does in deployment, the reported error rates can mislead users. After deployment, the detector produced false positives on human writing.

and some citations/further reading:

- Ralph B. D'Agostino, Scott Grundy, Lisa M. Sullivan, Peter Wilson, Validation of the Framingham Coronary Heart Disease prediction scores: results of a multiple ethnic groups investigation, JAMA, Volume 286, Issue 2, 2001, pp. 180–187, <https://jamanetwork.com/journals/jama/fullarticle/193997>
- Carrie J. Beede et al., A Human-Centered Evaluation of a Deep Learning System Deployed in Clinics for the Detection of Diabetic Retinopathy, CHI 2020, <https://doi.org/10.1145/3313831.3376718>
- C. Galen, R. Steele, Evaluating Performance Maintenance and Deterioration Over Time of Machine Learning-based Malware Detection Models on the EMBER PE Dataset, 2020 Seventh International Conference on Social Networks Analysis, Management and Security (SNAMS), Paris, France, 2020, pp. 1–7, <https://doi.org/10.1109/SNAMS52053.2020.9336538>

## Predictive feature concerns

Before modeling, ask:

- **Plausible signal:** Is there a defensible reason this should predict the target?
- **Artifacts and proxies:** Could the model use an accidental artifact?

Here are some examples from practice:

- **Plausible signal:** Researchers trained machine-learning models to infer personality traits from handwriting in a 2023 study and reported accuracies above 99%, despite controlled studies finding that graphology does not reliably measure personality (Dazzi & Pedrabissi, 2009; Neter & Ben-Shakhar, 1989). (The problem was that the researchers did not obtain independent personality measurements: they generated the personality labels from graphology rules based on the same handwriting features given to the model. The model therefore learned to reproduce those rules, not to predict actual personality.)
- **Artifacts and proxies:** During the COVID-19 pandemic, researchers trained chest X-ray models to distinguish COVID-19 pneumonia from other conditions. In some datasets, COVID-19 images came from different hospitals, scanners, or image acquisition protocols than the control images, allowing the model to learn hospital and equipment artifacts instead of disease related features. The models could score well on internal test sets that shared those artifacts but perform poorly on images from new hospitals. This failure is documented in [Common pitfalls and recommendations for using machine learning to detect and prognosticate for COVID-19](#).


and for citations/further reading, many more examples appear in:

- Mel Andrews, Andrew Smart, Abeba Birhane, The reanimation of pseudoscience in machine learning and its ethical repercussions, Nature Machine Intelligence, Volume 5, Issue 9, 2024, Article 101027, <https://doi.org/10.1016/j.patter.2024.101027>

**Example: distinguish real headshots from AI generated headshots.**

Suppose a career networking platform wants to enforce a “real photos only” policy. Engineers create a dataset from a **real face recognition dataset**, and fake headshots generated by an AI model.

The model achieves excellent test accuracy but fails in production. Looking at the data, we can see why:

| “Real” training image                                                             | Real deployment image                                                              |
|-----------------------------------------------------------------------------------|------------------------------------------------------------------------------------|
|  |  |

The “real” images in its training and test sets were cropped from candid photographs, while the generated images looked like posed studio headshots. The model could learn candid versus studio cues, including lighting, framing, and background, instead of real versus synthetic cues. This is an example of an “accidental artifact”.

This is a “Frankenstein dataset”: the label is correlated with the source dataset. Holding out samples from the same sources does not reveal the problem. In production, real profile photographs are also professional headshots, so the shortcut disappears.

| Earlier generator                                                                   | Newer generator                                                                      |
|-------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------|
|  |  |

Even after collecting real studio style headshots, the fake image distribution changes over time. A detector trained on images from an earlier generator may fail as newer generators produce more photorealistic portraits. This is an example of a “representativeness” concern.

The training data must reflect the model’s current deployment period. This dataset would need frequent updates and evaluation on images from generators not used for training.

## Other data quality concerns

- Label error
- Other data entry error
- Inconsistent units/formats
- Missing data

Some examples of data quality failures:

- You train a content classification model on a human-labeled dataset of 58K Reddit comments categorized according to 27 emotions. But up to 30% of the data is mislabeled, possibly because the annotators were shown text without context, did not understand slang or sarcasm, were unfamiliar with political or pop culture references, or just rushed through the job. For example, the sentence: “Really? Wow. You’re either hopelessly ignorant or you’re trolling. For your sake, I hope you’re trolling” is mislabeled as OPTIMISM.
- You train a ride hailing demand prediction model on GPS data from mobile phones. Some trips contain corrupted coordinates, such as a location recorded in the middle of the ocean (see: **Null Island**) or a misplaced decimal point showing a car moving 3,000 km in a minute. The model treats these entries as real trips, introducing noise and spurious patterns.
- You combine clinical datasets from multiple hospitals, but one site records lab results in mg/dL while another uses mmol/L. Your model makes nonsensical predictions.

Citations/further reading:

- Curtis G. Northcutt. Pervasive Label Errors in ML Datasets Destabilize Benchmarks. Blog post, March 29, 2021. <https://l7.curtisnorthcutt.com/label-errors>
- Curtis G. Northcutt, Anish Athalye, Jonas Mueller. Pervasive Label Errors in Test Sets Destabilize Machine Learning Benchmarks. In Advances in Neural Information Processing Systems (NeurIPS 2021). <https://neurips.cc/virtual/2021/47102>
- Surge AI, 30 Percent of Google’s Reddit Emotions Dataset Is Mislabeled. <https://surgehq.ai/blog/30-percent-of-googles-reddit-emotions-dataset-is-mislabeled>
- Brown, S. J., Goetzmann, W. N., Ibbotson, R. G., & Ross, S. A. Survivorship bias in performance studies. The Review of Financial Studies, 5(4), 553–580 (1992). <https://doi.org/10.1093/rfs/5.4.553>

## Stage 2: How do I process the data?

- Explore data, make and check assumptions
- Split data (avoid data leakage)
- Select features, target
- Construct samples
- Handle missing data
- Convert to numeric types
- Create “transformed” features

### Make and check assumptions

It’s always a good idea to “sanity check” your data. Before you look at it, think about what you expect to see. Review the data dictionary, and any documentation related to how the dataset was compiled.

Then, check to make sure your expectations are realized.

- Look at individual samples - aggregate trends can hide important details!
- Generate and look at plots of data distributions
- Compute summary statistics
- and consider general trends

## Split data

If we split data incorrectly, we can introduce *data leakage*, which causes an *overly optimistic evaluation*. Data leakage happens when information that would not be available in the real deployment setting “leaks” into the training or evaluation process, and makes the evaluation much “easier” than the real deployment task. We said earlier how important it is to avoid that!

The basic idea behind data splitting is the relationship between the training set and the evaluation set, should be the same as the relationship between the training set and the data we will see in deployment.

- if there is no structure in the data - shuffle split (avoid accidental patterns)
- if there is group structure - use a split that keeps members of each group in either training set, or validation set, but not both
- for time series data - use a split that keeps validation data in the future, relative to training data

In the following example, if we put the first samples in the training set and the last samples in the test set, we would accidentally create a distribution shift between training and test data, since there seems to be some order effect. For this data, we should randomly select samples to put in the test data - this is called a shuffle split.




Figure 2: This data requires a shuffle split.

In the next example, suppose we have multiple samples from the same individual.




Figure 3: This data has group structure.

In this case, the correct approach depends on how the model will be deployed:

- if the model will be asked to make predictions on new samples from the same people as in the training set, then the shuffle split is still OK! Here the evaluation task is to predict new samples from the people in the training set, which is the same as the deployment task.
- if, however, the model will be asked to make predictions for *new people* not in the training set, then the shuffle split will cause an overly optimistic evaluation. Predicting new samples from the people in the training set is an *easier* task than the deployment task! In this case, we must split the data so that each individual person is either in the training set, or in the test set.

In the next example, suppose we have a time sequence, and we want to predict a value that changes gradually over time, given the previous sequence:




Figure 4: Time series data.

In this case, a shuffle split (like the example on the left) will cause data leakage/an overly optimistic evaluation, because this makes the evaluation task much easier than the deployment task! Instead, we need to split the data chronologically (like the example on the right) to mimic the deployment task.




Figure 5: This data needs a chronological split.

## Select features

Include features if they are:

- Plausibly predictive
- Available at inference time
- Don't have other major issues

We must *exclude* fields in the data that have no plausible relationship to the target variable:

- A machine learning model will find “patterns” even if the feature data is not really related to the target variable! It will find “spurious” relationships. That can potentially be much worse than if there was no ML model at all.
- In many cases, there will be fields in the data that you shouldn't use, e.g. you won't use someone's phone number or home address to predict their Intro ML course grade.
- Always exclude fields that are purely identifiers (e.g. a numeric identifier for a sample).

We also must *exclude* fields in the data that are not going to be available to our model when it is deployed for “real”:

- Example: You build a model to help decide which applicants should be approved for loans. In the training dataset, you accidentally include fields such as “number of late payments,” which are only known after the loan has already been issued. By training on them, the model looks impressively accurate during evaluation, but in deployment those features are unavailable, so the model cannot make real world predictions.

Including fields that won't be available in deployment is also a type of data leakage/causes an overly optimistic evaluation. The evaluation task, where the extra feature is available, is “easier” than the deployment task where the feature is not available.

## Select target

Target should be:

- measureable
- available
- correct

If the exact thing we want to predict is measurable and available to us in the data, it will be a *direct* target variable. Sometimes, however, the thing we want to predict is not measurable or available. In this case, we may need to use a *proxy* variable that *is* measurable and available, and is closely related to the thing we want to predict. (The results will only be as good as the relationship between the thing we want to predict, and the proxy!)

Example: You want to predict a patient's illness severity in order to allocate extra care resources. But severity itself is not directly measurable in your dataset. Instead, you use past healthcare costs as a proxy variable, assuming that sicker patients generate higher costs. (In practice, this systematically underestimates the needs of disadvantaged groups who historically have had less access to care. The model appears accurate against the proxy, but fails against the true target.)

Citation: Ziad Obermeyer, Brian Powers, Christine Vogeli, Sendhil Mullainathan, Dissecting racial bias in an algorithm used to manage the health of populations, *Science*, Volume 366, Issue 6464, 2019, pp. 447–453, <https://doi.org/10.1126/science.aax2342>

We mentioned label error when we talked about data quality. Since it is often expensive to get labeled data, it's not uncommon for the target variable to be machine generated, or added by humans who spend very little time on each sample, and subject to high error rates.

## Construct samples

For each sample, you need:

- all features you plan to use
- the corresponding target value
- a reliable way to match them

Before training a supervised model, define what one sample represents: one person, image, transaction, event, or other unit of prediction.

Every feature and the target variable must be available at that individual sample level.

You may combine multiple data sources to “enrich” each sample. To do so, you need some key that establishes a one-to-one correspondence between the records. The feature values and target must describe the same person, object, or event.

If you cannot match the features to the target for each sample, you cannot construct the labeled dataset needed for supervised learning.

## Handle missing data

Missing data can appear in several different ways:

- Rows that have NaN values
- Rows that have other values encoding “missing” (-1, 0, 100...)
- Rows that are *not there* but should be

The first type is the easiest to detect!

The second type can be tricky if the value that is supposed to mean “missing” is also a feasible “real” value. For example, in the classic **NYC Taxi Dataset**, the “tip” field is recorded as a 0 for all trips that were paid in cash, because that data is not available for cash trips. However, it’s also plausible for the “tip” field to really be 0 - so if you didn’t look at the data dictionary and find out that cash tips aren’t recorded, you might think those were “real” 0s and not “missing” 0s.

The third type of missing data is easy to overlook because the missing values do not appear as blank cells—they appear as rows that are completely absent. For example, in a **traffic dataset**, a row might only be recorded when at least one vehicle travels from a particular entrance to a particular exit during a 15-minute period. If no vehicles make that trip, no row is created. However, that missing row actually contains useful information: the traffic count was zero. To accurately describe or model traffic, these missing rows should be added and recorded with a value of 0.

Missing data can occur for different reasons, and the cause of the missingness matters because it can affect how we analyze the data. It is often characterized as:

- **Missing Completely at Random (MCAR):** the chance that a value is missing does not depend on any observed or unobserved information. In other words, all samples are roughly equally likely to have missing values.
- **Missing at Random (MAR):** the chance that a value is missing depends on other information that we have observed. For example, whether income is missing might depend on a person’s age or occupation, both of which are recorded in the dataset.
- **Missing Not at Random (MNAR):** the chance that a value is missing depends on the missing value itself. For example, people with very high incomes might be more likely to leave the income question blank.

How should you handle small amounts of missing data?

- You may leave the value as missing (e.g. NaN) if your modeling or analysis method can handle missing values directly.
- You may omit the row if the amount of missing data is small and the data is Missing Completely at Random (MCAR), so removing those rows is unlikely to introduce bias. You may also consider dropping a column if it contains too much missing data or is not important to the analysis.
- Otherwise, you may impute (fill in) the missing values, with the goal of estimating the unknown value as reasonably as possible.

Some simple imputation methods include:

- filling with the mean, median, or mode
- using forward fill or interpolation for ordered or time-series data, when nearby observations are likely to be informative

### Convert to numeric types

- ordinal and one-hot encoding of categorical data
- text to “bag of words” or other representation
- image data to raw pixels
- audio to frequency domain (or image of frequency domain) features

For most machine learning models, we will need all data to be represented using some kind of numeric data type!

Suppose we have a fragrance dataset, and we want to prepare to train a model to predict rating for new perfumes:

| Perfume       | Perfumer  | Longevity | Family | Description            | Rating |
|---------------|-----------|-----------|--------|------------------------|--------|
| Rose Veil     | A. Rivera | short     | floral | soft rose with pepper  | 4.2    |
| Blooming Iris | A. Rivera | medium    | floral | soft dry iris          | 4.8    |
| Pepper Ember  | J. Kim    | medium    | spicy  | warm pepper            | 4.6    |
| Saffron Smoke | J. Kim    | long      | spicy  | warm saffron and cedar | 4.0    |
| Cedar Night   | L. Patel  | long      | woody  | dry cedar              | 4.4    |
| Forest Moss   | L. Patel  | short     | woody  | dry moss               | 4.7    |

Here, the longevity, fragrance family, name of the perfumer, and the description may all relate to the rating, but they are not numeric. We need to transform them into some type of numeric format.

The Longevity field and the Family field are both categorical: they can take on a discrete set of possible values. However, the Longevity field has a natural order (short < medium < long), while there is no logical order to floral, spice, woody.

For an ordinal categorical variable, where the categories have a meaningful order, we can use ordinal encoding, replacing each category with a number that preserves the ordering. For example, we could encode Longevity as:

| Longevity | Encoded value |
|-----------|---------------|
| short     | 0             |
| medium    | 1             |
| long      | 2             |

The transformed `Longevity` column would therefore look like:

| Perfume       | Longevity | Longevity_encoded |
|---------------|-----------|-------------------|
| Rose Veil     | short     | 0                 |
| Blooming Iris | medium    | 1                 |
| Pepper Ember  | medium    | 1                 |
| Saffron Smoke | long      | 2                 |
| Cedar Night   | long      | 2                 |
| Forest Moss   | short     | 0                 |

This encoding makes sense because the numbers preserve the natural ordering: `short < medium < long`.

For a nominal categorical variable, where the categories do not have a meaningful order, we generally do not want to assign arbitrary numbers. For example, encoding `floral = 0`, `spicy = 1`, and `woody = 2` could incorrectly suggest that `woody` is somehow “larger than” `spicy`, or that `spicy` lies between `floral` and `woody`.

Instead, we can use one-hot encoding. One-hot encoding creates a separate binary column for each category. A value of 1 indicates that the sample belongs to that category, while 0 indicates that it does not.

For `Family`, the transformation would be:

| Perfume       | Family | Family_floral | Family_spicy | Family_woody |
|---------------|--------|---------------|--------------|--------------|
| Rose Veil     | floral | 1             | 0            | 0            |
| Blooming Iris | floral | 1             | 0            | 0            |
| Pepper Ember  | spicy  | 0             | 1            | 0            |
| Saffron Smoke | spicy  | 0             | 1            | 0            |
| Cedar Night   | woody  | 0             | 0            | 1            |
| Forest Moss   | woody  | 0             | 0            | 1            |

The `Perfumer` field could be one-hot encoded, but what if the set of possible perfumers is very large? Then, one-hot encoding would create a large number of additional, mostly sparse columns.

Instead, we could use target encoding. With target encoding, each category is replaced by a numeric value based on the target variable.

For example, if our target is `Rating`, we could replace each perfumer with the average rating of perfumes made by that perfumer:

| Perfumer  | Ratings  | Target-encoded value |
|-----------|----------|----------------------|
| A. Rivera | 4.2, 4.8 | 4.50                 |
| J. Kim    | 4.6, 4.0 | 4.30                 |
| L. Patel  | 4.4, 4.7 | 4.55                 |

The transformed column would then look like:

| Perfume       | Perfumer  | Perfumer_encoded |
|---------------|-----------|------------------|
| Rose Veil     | A. Rivera | 4.50             |
| Blooming Iris | A. Rivera | 4.50             |
| Pepper Ember  | J. Kim    | 4.30             |
| Saffron Smoke | J. Kim    | 4.30             |
| Cedar Night   | L. Patel  | 4.55             |
| Forest Moss   | L. Patel  | 4.55             |

This reduces a potentially large categorical variable to a single numeric column. However, target encoding must be used carefully because it uses information from the target variable itself. If the encoding is calculated using the full dataset, information from the test data can leak into the training data. The encoding should be learned only from the training data.

What happens if a perfumer appears in the test set that never appeared in the training set? Since we do not have a target-encoded value for that perfumer, a common approach is to use the overall average target value from the training set. In this example, the average training-set rating is approximately 4.45, so a new perfumer could be encoded as:

| Perfume    | Perfumer | Perfumer_encoded |
|------------|----------|------------------|
| Amber Tide | M. Chen  | 4.45             |

This gives the model a reasonable default value.

The `Description` field contains free-form text, not categories, so we need a different way to turn it into numbers. One simple approach is *bag of words*. Bag of words creates a vocabulary containing the words that appear in the training descriptions, then creates one column for each word in the vocabulary. Each description is represented by the number of times each vocabulary word appears.

For example, suppose we use the vocabulary:

soft, rose, iris, warm, pepper, saffron, dry, cedar, moss

(omitting words like `and` and `with`, which are not relevant to the task).

Then the descriptions could be transformed like this:

| Perfume       | soft | rose | iris | warm | pepper | saffron | dry | cedar | moss |
|---------------|------|------|------|------|--------|---------|-----|-------|------|
| Rose Veil     | 1    | 1    | 0    | 0    | 1      | 0       | 0   | 0     | 0    |
| Blooming Iris | 1    | 0    | 1    | 0    | 0      | 0       | 1   | 0     | 0    |
| Pepper Ember  | 0    | 0    | 0    | 1    | 1      | 0       | 0   | 0     | 0    |
| Saffron Smoke | 0    | 0    | 0    | 1    | 0      | 1       | 0   | 1     | 0    |
| Cedar Night   | 0    | 0    | 0    | 0    | 0      | 0       | 1   | 1     | 0    |
| Forest Moss   | 0    | 0    | 0    | 0    | 0      | 0       | 1   | 0     | 1    |

Each text description is now represented by a set of numeric features that can be used by a model. In a real dataset, the vocabulary would usually contain many more words, so this transformation can create a large and sparse set of columns.

Bag of words ignores word order and context. It is therefore a simple representation of which words appear in a description, rather than a representation of the full meaning of the text.

As with other preprocessing steps, the vocabulary should be learned from the training set only. If a new description in the test set contains a word that was never seen during training, that word is simply not represented by one of the existing bag-of-words columns.

Other types of data also need to be converted into numeric features before they can be used by most models.

- For image data, one simple representation is the raw pixel values: each pixel is represented by numeric intensity or color values, such as RGB (red, green, blue) values.
- For audio data, the raw waveform can be transformed into the frequency domain, which describes how much energy is present at different frequencies. These frequency-based features are often represented using a spectrogram, which is essentially an image showing how the frequency content of the audio changes over time. The resulting numeric values, or the spectrogram itself, can then be used as features for a model.

### Create “transformed” features

Sometimes the raw features are not in the most useful form for a model, so we create new features from them. For example,

- instead of using a timestamp directly, we might create an `is_rush_hour` feature that indicates whether the observation occurred during a typical commuting period.
- we might extract `hour_of_day`, `day_of_week`, `month`, or `is_weekend` from a date
- we might compute `age` from a date of birth, `price_per_square_foot` from price and area, `distance_from_city_center` from latitude and longitude
- we might create a `days_since_last_purchase` feature from transaction dates
- we might take the log of a highly skewed variable such as income
- we might combine measurements into a ratio such as `debt_to_income`.

The goal is to represent the raw data in a form that makes patterns easier for the model to learn. This will be highly data- and context-specific!

### Recap: Working with data

This lesson focused on the work that must happen before and around model training and evaluation.

First, ask whether the data is appropriate: was it collected legally and ethically, does it represent the deployment setting, does it contain plausible predictive signals rather than accidental artifacts, and is it accurate and complete?

Then prepare the data for modeling. Define what each sample represents, select features and a target, split the data without leakage, handle missing values, convert non-numeric data into numeric features, and create useful transformed features.

Data is also central to evaluation: if the evaluation data is not representative, or if information from it leaks into training, the measured performance can be overly optimistic and fail to predict how the model will perform in deployment. This is a much worse outcome than simply having a poor-performing model (but knowing it is poor)!