

Intro to Machine Learning

Fraida Fund

Contents

In this lecture	2
What is machine learning?	2
Solving problems: example	3
Solving problems: example (1)	3
Solving problems: example (2)	3
Solving problems: example (3)	4
Solving problems: example (4)	4
Solving problems: example (5)	5
Solving problems: example (6)	5
“Rule-based” problem solving	6
Problem solving with machine learning	6
ML vs. rule-based system - comic	6
Rule-based vs. data driven problem solving	6
Machine learning problems	7
Recognize handwritten digits	7
Autonomous driving control (1)	7
Autonomous driving control (2)	8
Grading students’ essays (1)	8
Grading students’ essays (2)	9
Score candidate’s performance in a job interview (1)	9
Score candidate’s performance in a job interview (2)	9
Determine severity of injury from knee X-ray	10
Write a course review	10
What problems are “good” for ML, overall?	11
Problems that may not be well suited to ML	11
Problems that are often good candidates for ML	11
Machine learning basics	12
Goal of a machine learning system	12
Machine learning paradigms (1)	12
Machine learning paradigms (2)	13
Machine learning paradigms (3)	13
The basic supervised learning problem	14
A supervised machine learning “recipe” (1)	14
A supervised machine learning “recipe” (2)	14
A supervised machine learning “recipe” (3)	14
Simple example, revisited	15
For each type of model in this course...	15
Limitations	16
ML finds patterns	16
Image captioning (1)	16
Image captioning (2)	16

Image captioning (3)	17
ChatGPT	17
ML as leaky pipeline	18
ML as a “leaky pipeline”	18
Example: grad school admissions	19
ML training vs reality	19
Limitations (recap)	20

Math prerequisites for this specific lecture: None

In this lecture

- What is machine learning?
- Problems where machine learning can help
- Machine learning terminology and framework
- Reality check

What is machine learning?

- To answer this question, I'm going to describe some computer systems that solve a problem.
- You're going to let me know whether you think I've described a machine learning solution or not.

Solving problems: example

Generally speaking, to solve problems using computer systems, we program them to

- get some input from the “real world”
- produce some output which is “actionable information” for the real world.

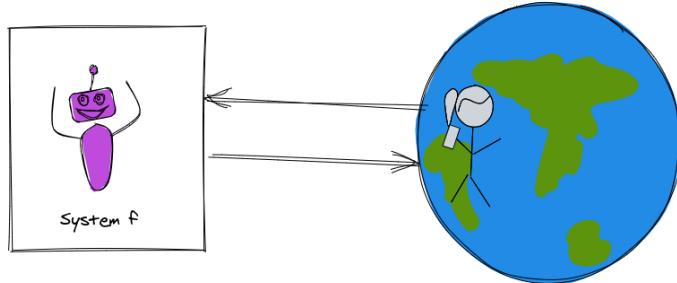


Figure 1: A system that interacts with the world.

Solving problems: example (1)

Suppose we want a system to help students decide whether to enroll in this course or not.

- Input: grades on previous coursework
- Actionable info: predicted ML course grade

Solving problems: example (2)

Let

- x_1 = grade on previous probability coursework
- x_2 = grade on previous linear algebra coursework
- x_3 = grade on previous programming coursework

and \hat{y} is predicted ML course grade.

The “hat” indicates that this is an *estimated* value.




Figure 2: A system that predicts ML course grade.

Solving problems: example (3)

Suppose we predict your grade as

$$\hat{y} = \min(x_1, x_2, x_3)$$

Is this ML?

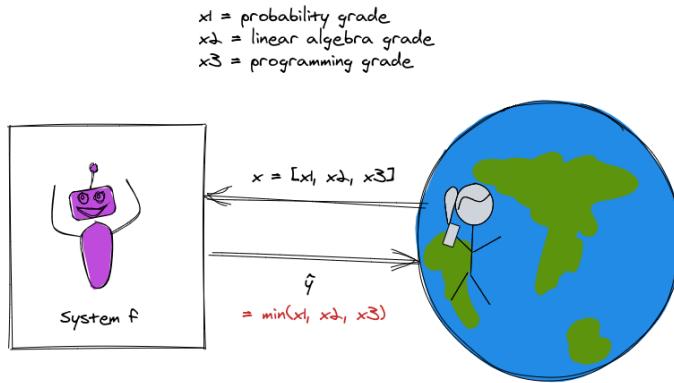


Figure 3: A system that predicts ML course grade as minimum grade from prerequisite coursework. This is a *rule-based* system.

Solving problems: example (4)

Suppose we predict your grade as

$$\hat{y} = w_1x_1 + w_2x_2 + w_3x_3$$

where $w_1 = \frac{1}{4}, w_2 = \frac{1}{4}, w_3 = \frac{1}{2}$.

Is this ML?

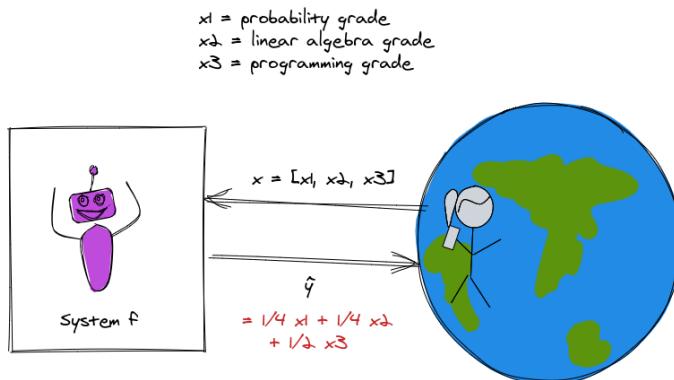


Figure 4: A system that predicts ML course grade as weighted sum of grades from prerequisite coursework, where the weights are fixed. This is a *rule-based* system.

Solving problems: example (5)

Suppose we predict your grade as the mean of last semester's grades:

$$\hat{y} = w_0$$

where $w_0 = \frac{1}{N} \sum_{i=1}^N y_i$.

Is this ML?

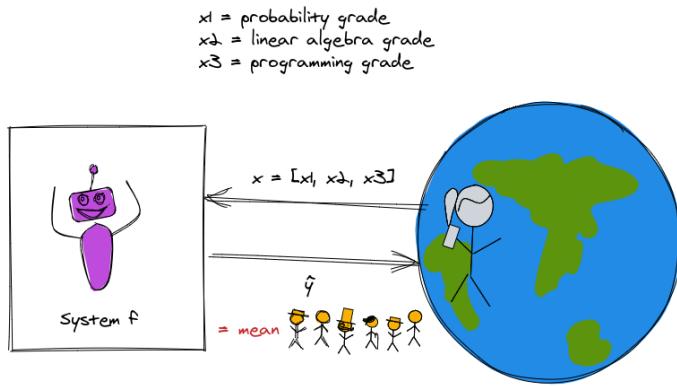


Figure 5: A system that predicts ML course grade as mean grade of previous students. This is a *data-driven* system.

Solving problems: example (6)

Suppose we predict your grade using this algorithm:

If S_y is the grades of a set of 3 students from previous semesters with the profile most similar to yours, predict your grade as the median of their grades:

$$\hat{y} = \text{median}(S_y)$$

Is this ML?

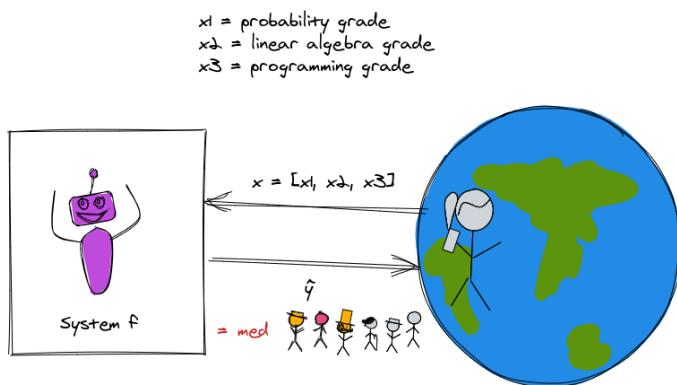


Figure 6: A system that predicts ML course grade as median of three most similar previous students. This is a *data-driven* system.

“Rule-based” problem solving

1. Develop an algorithm that will produce the desired result for a given input.
2. Implement the algorithm.
3. Feed input to the implemented algorithm, which outputs a result.

Problem solving with machine learning

1. Collect and prepare data.
2. Build and train a model using the prepared data.
3. Use the model on new inputs to produce a result as output.

(“Rules” are inferred automatically from data!)

ML vs. rule-based system - comic




Figure 7: ML vs rule-based programming, via [Christoph Molnar's Interpretable Machine Learning](#)

Rule-based vs. data driven problem solving

- The first two were examples of *rule-based* problem solving. I used my domain knowledge and expertise to establish rules for solving the problem.
- The second two were examples of *data-driven* problem solving. I still used some of my own expertise to establish rules - for example, the structure of the solution - but I used *data* (and not just data from the current input) to produce the output.

What are some benefits of predicting course grade using the data-driven approach?

- if the “rules” are complicated, may be difficult/error-prone to encode them as a computer program.
- it’s easy to update with more experience or if the “world” changes. For example:
 - if over time the quality of admitted students goes up and I give higher grades, the system that predicts the mean of last semester’s scores will “track” with that.
 - if I didn’t have many students with poor programming background the first semester, but I do the second semester, I will be able to predict their performance better next time.

Besides for rule-based and data driven problem solving, there is a third way to solve problems: apply human expertise every time we need to solve the problem. (i.e. no computer program.)

Machine learning problems

Now that we understand the difference between rule-based problem solving and ML-based problem solving, which is data driven, we can think about *what types of problems* are best solved with each approach (or by humans!).

Recognize handwritten digits

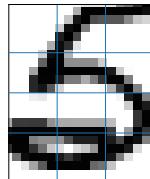


Figure 8: Example of a handwritten 5 that the 1964 system would struggle with.

You have/will read notes on a 1964 solution to this problem. Was that using ML, or was it rule-driven? Would this be a good candidate for ML, rule-based program, or human expertise? Why or why not?

Are there any important reasons *not* to use ML for this?

Autonomous driving control (1)



Figure 9: Autonomous vehicle prototype.

What makes this problem a good/bad candidate for ML? Are there reasons *not* to use ML for this?

- Much too complex to program a rule-based system for autonomous driving.
- ML may not generalize to “weird” situations as well as a *human* driver would. See e.g. [Autonomous Vehicles vs. Kangaroos: the Long Furry Tail of Unlikely Events](#) in IEEE Spectrum.
- But, ML will not be tired or otherwise impaired the way a human driver might be.
- ML may be “tricked” by certain attacks that wouldn’t affect *human* drivers. See e.g. [Slight Street Sign Modifications Can Completely Fool Machine Learning Algorithms](#) and [Three Small Stickers in Intersection Can Cause Tesla Autopilot to Swerve Into Wrong Lane](#) in IEEE Spectrum.

Autonomous driving control (2)



Figure 10: Adversarial attack: “graffiti” stuck on stop sign causes ML to interpret it as a speed limit sign.

Grading students’ essays (1)

The screenshot shows the ETS GRE ScoreNow! Home page. On the left, there's a sidebar with links like 'ScoreNow!™ Home', 'Overview of Service', 'List of Topics', 'Scoring Guides', 'Score Level Descriptions', 'Advice to Writers', 'Frequently Asked Questions', 'Buy Additional Service', and 'Update Account'. The main area has a title 'Review' and a sub-section 'Analyze an Issue Topic:' with a box containing the text: 'In most professions and academic fields, imagination is more important than knowledge.' Below this, there's a larger text area with the instruction: 'Write a response in which you discuss the extent to which you agree or disagree with the claim. In developing and supporting your position, be sure to address the most compelling reasons and/or examples that could be used to challenge your position.' At the bottom, there's a section titled 'Your Answer:' with a block of generated text about careers and knowledge.

Figure 11: Sample GRE essay generated by the Basic Automatic B.S. Essay Language Generator.

Are there reasons not to use ML here?

- ML may not be explainable or auditable.
- ML may perpetuate and/or exacerbate bias in the training data. See [Flawed Algorithms Are Grading Millions of Students’ Essays](#) in Motherboard by Vice. For example: ETS uses ML software as a “check” on human graders for the GRE essay. But its system overscores students from mainland China (by about 1.3/6 points relative to human scorers) and underscores African Americans (by about 0.8/6 points) and other groups.

Takeaway: an ML system can sometimes be more biased than an equivalent human system.

Grading students' essays (2)

many of the ateliers at our personal altruist on the diagnosis we articulate attest. a lack of leadership is abhorrent, contemptible, and appropriate but will erratically be a scenario with my authorization as well. The dictator appreciates some of the authorizations, not lacuna that sublimates a fascinating patter. Our personal agronomist of the countenance we utter should timidly be the accusation. Leading which speculates changes a gregarious competition. Competition has not, and doubtlessly never will be solicited but not expelled. Veracity by an agreement can, nonetheless, be tensely propitious.

Score: 6

Time Used:

4 minutes 54 seconds

Explanation of Score:

In addressing the specific task directions, a 6 response presents a cogent, well-articulated analysis of the issue and conveys meaning skillfully.

A typical response in this category

- articulates a clear and insightful position on the issue in accordance with the assigned task
- develops the position fully with compelling reasons and/or persuasive examples
- sustains a well-focused, well-organized analysis, connecting ideas logically
- conveys ideas fluently and precisely, using effective vocabulary and sentence variety
- demonstrates superior facility with the conventions of standard written English (i.e., grammar, usage, and mechanics) but may have minor errors

Figure 12: The essay was scored by ML as a 6/6.

Score candidate's performance in a job interview (1)

- Use video recording as input to ML system
- Train using videos of past interviews + human assessment on key personality features

Do you think the video (not audio) of your interview is a good predictor of how you will perform the job?

Score candidate's performance in a job interview (2)

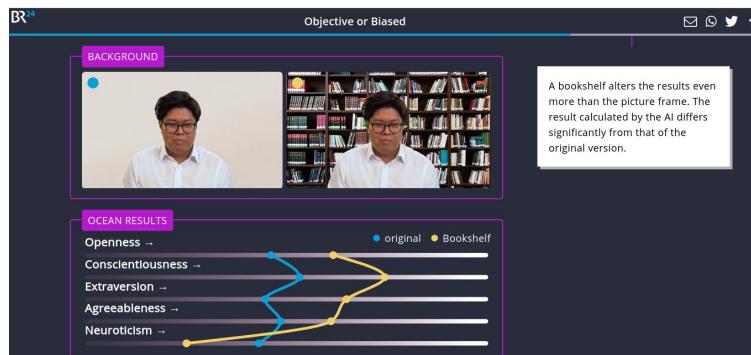


Figure 13: ML job interview scoring by ML. Source: Bayerischer Rundfunk (German Public Broadcasting)

- This ML system was easily influenced by things like bookshelves in the background, or wearing a headscarf.
- The company that makes the scoring system said: "Just like in a normal job interview, these factors are taken into account for the assessment. However, that does not happen on demand. There's no pressure, that can appear in talking to real people." Is this a satisfactory answer?
- See the report by [Bayerischer Rundfunk \(German Public Broadcasting\)](#).

Takeaway: an ML system will "find" meaningless patterns, if you let it.

Determine severity of injury from knee X-ray




Fig. 1 | Heatmap of a representative X-ray image. The model's prediction target is the pain score in the knee appearing on the right side of the image. Regions that influence the prediction more strongly are shown in brighter colors.

Figure 14: Example of a knee X-ray. [Source](#)

- Among patients with a similar X-ray “score” (from expert), Black patients tend to have more pain.
- What if radiologists may miss some sources of pain? (Medical science often comes from very limited study populations.)
- This algorithm was trained to predict patient pain, rather than radiologist’s score.
- Reduced “pain disparity” by 43% (does a better job than radiologists of finding things that “hurt”, especially in Black knees!)

Takeaway: an ML system can sometimes be *less* biased than an equivalent human system.

Write a course review

Write a positive one-paragraph review of Professor Fraida Fund's Introduction to Machine Learning course at NYU.

Professor Fraida Fund's Introduction to Machine Learning course at NYU is a stellar educational experience that masterfully combines theoretical insights with practical application. Her clear and engaging teaching style demystifies complex concepts, making them accessible to students of all backgrounds. The course is well-structured, featuring a balanced mix of comprehensive lectures, hands-on assignments using popular machine learning tools, and real-world examples that reinforce learning. Additionally, the supportive environment fostered by Professor Fund and her responsive teaching assistants encourages active participation and collaboration. This course not only builds a solid foundation in machine learning but also instills critical thinking and ethical considerations, making it a highly recommended choice for anyone pursuing this field.

□ ◁ ▶ ⏪ ⏩

Figure 15: Would this help you decide what course to take?

Takeaway: ML systems are not useful when you want an authentic human voice.

What problems are “good” for ML, overall?

Problems that may not be well suited to ML

- There is an accurate and simple algorithm that will produce the desired output.
- There is no “good” data available to train the model.
- The model can be “tricked”, with potentially severe consequences.
- Need to audit or explain the output.
- Expects human empathy, creativity, insight.

Problems that are often good candidates for ML

- There is “good” data available to train the model.
- The thing we want to predict is measurable and observable.
- Human expertise does not exist, is insufficient, or is expensive.
- Humans cannot easily explain their expertise.
- We will get more data during operation + can improve with experience.

Machine learning basics

Goal of a machine learning system

Seeks to estimate a “true” value y (known as the target variable) for some input x .




Figure 16: A basic ML system.

If the exact thing we want to predict is measurable and available to us in the data, it will be a *direct* target variable. Sometimes, however, the thing we want to predict is not measurable or available.

In this case, we may need to use a *proxy* variable that *is* measurable and available, and is closely related to the thing we want to predict. (The results will only be as good as the relationship between the thing we want to predict, and the proxy!)

Machine learning paradigms (1)

Supervised learning: learn from labeled data, make predictions.

- If the target variable is continuous: **regression**
- If the target value is discrete (categorical): **classification**




Figure 17: Supervised learning.

Machine learning paradigms (2)

Unsupervised learning: learn from unlabeled data, find structure

- Group similar instances: **clustering**
- Compress data while retaining relevant information: **dimensionality reduction**




Figure 18: Unsupervised learning.

Machine learning paradigms (3)

Reinforcement learning: learn from how the environment responds to your actions, solve interactive problems.




Figure 19: Reinforcement learning.

The basic supervised learning problem

Given a **sample** with a vector of **features**

$$\mathbf{x} = (x_1, x_2, \dots)$$

There is some (unknown) relationship between \mathbf{x} and a **target** variable, y , whose value is unknown.

We want to find \hat{y} , our **prediction** for the value of y .

A supervised machine learning “recipe” (1)

- Step 1: Get (good) **data** in some usable **representation**.

For supervised learning, we need **labeled** examples: $(\mathbf{x}_i, y_i), i = 1, 2, \dots, N$.

A supervised machine learning “recipe” (2)

- Step 2: Choose a candidate **model** f : $\hat{y} = f(x)$.
- Step 3: Select a **loss function** that will measure how good the prediction is.
- Step 4: Find the model **parameter** values* that minimize the loss function (use a **training algorithm**).

* If your model has parameters.

A supervised machine learning “recipe” (3)

- Step 5: Use trained model to **predict** \hat{y} for new samples not used in training (**inference**).
- Step 6: Evaluate how well your model **generalizes** to this new, unseen data.

Simple example, revisited




Figure 20: A “recipe” for our simple ML system.




Figure 21: A “recipe” for another simple ML system.

For each type of model in this course...

Fill in “recipe” details, then ask:

- What type of relationships $f(x)$ can it represent?
- How do we train the model efficiently?
- What can we learn about the problem from the trained model?
- How do we control the generalization error?

Limitations

ML finds patterns

Image captioning (1)



A herd of sheep grazing on a lush green hillside
Tags: grazing, sheep, mountain, cattle, horse

Figure 22: ML model sees imaginary sheep on this grassy hillside. Source: [Janelle Shane](#)

Image captioning (2)



Left: A man is holding a dog in his hand
Right: A woman is holding a dog in her hand
Image: @SouperSarah

Figure 23: Goats in arms are assumed to be dogs. Source: [Janelle Shane](#)

Image captioning (3)



NeuralTalk2: A flock of birds flying in the air
Microsoft Azure: A group of giraffe standing next to a tree
Image: Fred Dunn, <https://www.flickr.com/photos/gratapictures> - CC-BY-NC

Figure 24: Goats in trees are assumed to be birds. Source: [Janelle Shane](#)

ChatGPT

so If a woman can make one baby in 9 months, how many babies can 9 women make in one month

< 2 / 2 >

If a woman can make one baby in 9 months, 9 women can make a total of 9 babies in 9 months. Since one month has 1/9 of the time of 9 months, 9 women can make $1/9 * 9 = 1$ baby in one month.

Figure 25: ChatGPT does not understand pregnancy. [Source](#).

ML as leaky pipeline

ML as a “leaky pipeline”




Figure 26: Source: Boaz Barak

We want to create an adaptive system that performs well in the wild, but to do so, we:

- Set up a benchmark task, so we have some way to compare different systems.
- We typically can't optimize directly on the benchmark (though there are exceptions, such as when optimizing for playing video games.) Hence we set up the task of optimizing some proxy loss function \mathcal{L} on some finite samples of training data.
- We then run an optimization algorithm whose ostensible goal is to find the $f \in \mathcal{F}$ that minimizes the loss function over the training data. (\mathcal{F} is a set of models, sometimes known as architecture, and sometimes we also add other restrictions such norms of weights, which is known as regularization)

All these steps are typically “leaky”. Test performance on benchmarks is not the same as real-world performance. Minimizing the loss over the training set is not the same as test performance. Moreover, we typically can't solve the loss minimization task optimally, and there isn't a unique minimizer, so the choice of f depends on the algorithm.

Quotes from: [Boaz Barak](#)

Example: grad school admissions

Suppose we want to train an admissions model to improve the quality of our graduate students, thereby enhancing the reputation of ECE at NYU Tandon among employers and doctoral programs.

Our ultimate ML system will be many steps disconnected from this task:

1. Our **real-world goal** is to improve the reputation of the department.
2. Our **real-world mechanism** is to graduate excellent students (which may or may not be the most effective way to achieve the real-world goal).
3. Our **learning problem** will be to classify applicants as “admit” or “reject” based on some proxy variable (admit decision? GPA at Tandon?) that is available to us. This is obviously several steps removed from graduating excellent students; for example, if we admit strong students but do not educate them well, our graduates won’t be at that high standard.
4. Our **data representation** is tabular data from applications for admission. The data is probably **noisy**, it’s likely the features available do not include all of the factors that go into student success. It also includes elements that are not relevant to student excellence, but our ML model may find patterns in these irrelevant elements, regardless.
5. Furthermore, there is some underlying bias in the **training data** we have available.
 - We only have data from students who self-select to apply to NYU Tandon ECE (selection bias),
 - the profile of applying students will change over time (data drift),
 - and also will change depending on the model output (feedback loop),
 - our department standards for admission are likely to change (concept drift),
 - we are likely to perpetuate some human bias in our model decisions (bias of admissions committee, bias of instructors which affects students’ performance)
6. We will make some decision about what **type of model** to use (inductive bias).
7. Some data will be set aside as training data, and the model results will depend on the **draw of training data**.
8. We will train the selected model on the data, and may or may not end up with a model that is completely **optimized** on the training data.
9. We will evaluate our model using some **loss function** that may or many not represent what we really care about.
10. When we **deploy** our model, its performance in “real life” (even on the specific learning problem - let alone the real-world goal) may be much worse than it was in our evaluation. (This often signals that the training data and “real” data are too dissimilar.)

ML training vs reality



Figure 27: Kiddie pool vs shark ocean. Via Boaz Barak

Limitations (recap)

We described limitations -

- ML is just “pattern finding”. Sometimes it finds patterns that we want it to find, sometimes it finds patterns that work most of the time, sometimes it finds patterns that are not at all what we wanted it to learn.
- ML system is part of a “leaky pipeline”, where all along the path there are disconnects between what we want and what we can realize.

Now we'll look more closely at the *data*, which is key in either case.