

# Model selection

Fraida Fund

## Contents

Model selection problems . . . . .	2
Choosing model complexity . . . . .	2
Model order selection problem . . . . .	2
Using loss function for model order selection? . . . . .	2
Feature selection problem . . . . .	2
Validation . . . . .	3
Simple train/validation/test split . . . . .	3
Simple train/validation/test algorithm (1) . . . . .	3
Simple train/validation/test algorithm (1) . . . . .	3
Problems with simple split . . . . .	3
K-fold cross validation . . . . .	4
K-fold CV illustrated . . . . .	4
K-fold CV - algorithm (1) . . . . .	4
K-fold CV - algorithm (2) . . . . .	4
K-fold CV - how to split? . . . . .	6
One standard error rule . . . . .	7
One standard error rule - algorithm (1) . . . . .	7
One standard error rule - algorithm (2) . . . . .	7
One standard error rule - algorithm (3) . . . . .	7

## Model selection problems

Model selection problem: how to select the  $f()$  that maps features  $X$  to target  $y$ ?

We'll look at two examples of model selection problems, but there are many more.

### Choosing model complexity

We need to select a model of appropriate complexity -

- what does that mean, and
- how do we select one?

### Model order selection problem

- Given data  $(x_i, y_i), i = 1 \dots, N$  (one feature)
- Polynomial model:  $\hat{y} = w_0 + w_1x + \dots + w_dx^d$
- $d$  is degree of polynomial, called **model order**
- **Model order selection problem:** choosing  $d$

### Using loss function for model order selection?

Suppose we would "search" over each possible  $d$ :

- Fit model of order  $d$  on training data, get  $\mathbf{w}$
- Compute predictions on training data
- Compute loss function on training data:  $MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$
- Select  $d$  that minimizes loss
- Problem: loss function always decreasing with  $d$  (training error decreases with model complexity!)

### Feature selection problem

Given high dimensional data  $\mathbf{X} \in R^{n \times d}$  and target variable  $y$ ,

Select a subset of  $k \ll d$  features,  $\mathbf{X}_S \in R^{n \times k}$  that is most relevant to target  $y$ .

- Linear model:  $\hat{y} = w_0 + w_1x_1 + \dots + w_dx_d$
- Many features, only some are relevant
- **Feature selection problem:** fit a model with a small number of features

Why use a subset of features?

- High risk of overfitting if you use all features!
- For linear regression, there's a unique OLS solution only if  $n \geq d$
- For linear regression, when  $N \geq p$ , variance increases linearly with number of parameters, inversely with number of samples. (Not derived in class, but read extra notes posted after class at home.)

## Validation

### Simple train/validation/test split

- Divide data into training, validation, test sets
- For each candidate model, learn model parameters on training set
- Measure error for all models on validation set
- Select model that minimizes error on validation set
- Evaluate model on test set

Note: sometimes you'll hear "validation set" and "test set" used according to the reverse meanings.

### Simple train/validation/test algorithm (1)

- Split  $X, y$  into training, validation, and test.
- Loop over models of increasing complexity: For  $p = 1, \dots, p_{max}$ ,
  - **Fit:**  $\hat{w}_p = \text{fit}_p(X_{tr}, y_{tr})$
  - **Predict:**  $\hat{y}_{v,p} = \text{pred}(X_v, \hat{w}_p)$
  - **Score:**  $S_p = \text{score}(y_v, \hat{y}_{v,p})$

### Simple train/validation/test algorithm (1)

- Select model order with best score (here, assuming "lower is better"):

$$p^* = \underset{p}{\operatorname{argmin}} S_p$$

- Evaluate:

$$S_{p^*} = \text{score}(y_{ts}, \hat{y}_{ts,p^*}), \quad \hat{y}_{ts,p^*} = \text{pred}(X_{ts}, \hat{w}_{p^*})$$

### Problems with simple split

- Fitted model (and test error!) varies a lot depending on samples selected for training and validation.
- Fewer samples available for estimating parameters.
- Especially bad for problems with small number of samples.

## K-fold cross validation

Alternative to simple split:

- Divide data into  $K$  equal-sized parts (typically 5, 10)
- For each of the “splits”: evaluate model using  $K - 1$  parts for training, last part for validation
- Average the  $K$  validation scores and choose based on average

## K-fold CV illustrated

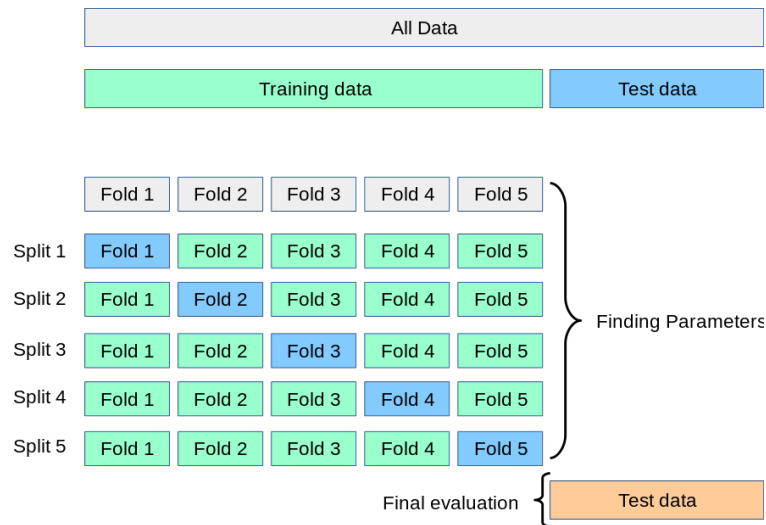


Figure 1: K-fold CV

## K-fold CV - algorithm (1)

**Outer loop** over folds: for  $i = 1$  to  $K$

- Get training and validation sets for fold  $i$ :
- **Inner loop** over models of increasing complexity: For  $p = 1$  to  $p_{max}$ ,
  - **Fit:**  $\hat{w}_{p,i} = \text{fit}_p(X_{tr,i}, y_{tr,i})$
  - **Predict:**  $\hat{y}_{v,i,p} = \text{pred}(X_{v,i}, \hat{w}_{p,i})$
  - **Score:**  $S_{p,i} = \text{score}(y_{v,i}, \hat{y}_{v,i,p})$

## K-fold CV - algorithm (2)

- Find average score (across  $K$  scores) for each model:  $\bar{S}_p$
- Select model with best *average* score:  $p^* = \text{argmin}_p \bar{S}_p$
- Re-train model on entire training set:  $\hat{w}_{p^*} = \text{fit}_{p^*}(X_{tr}, y_{tr})$
- Evaluate new fitted model on test set

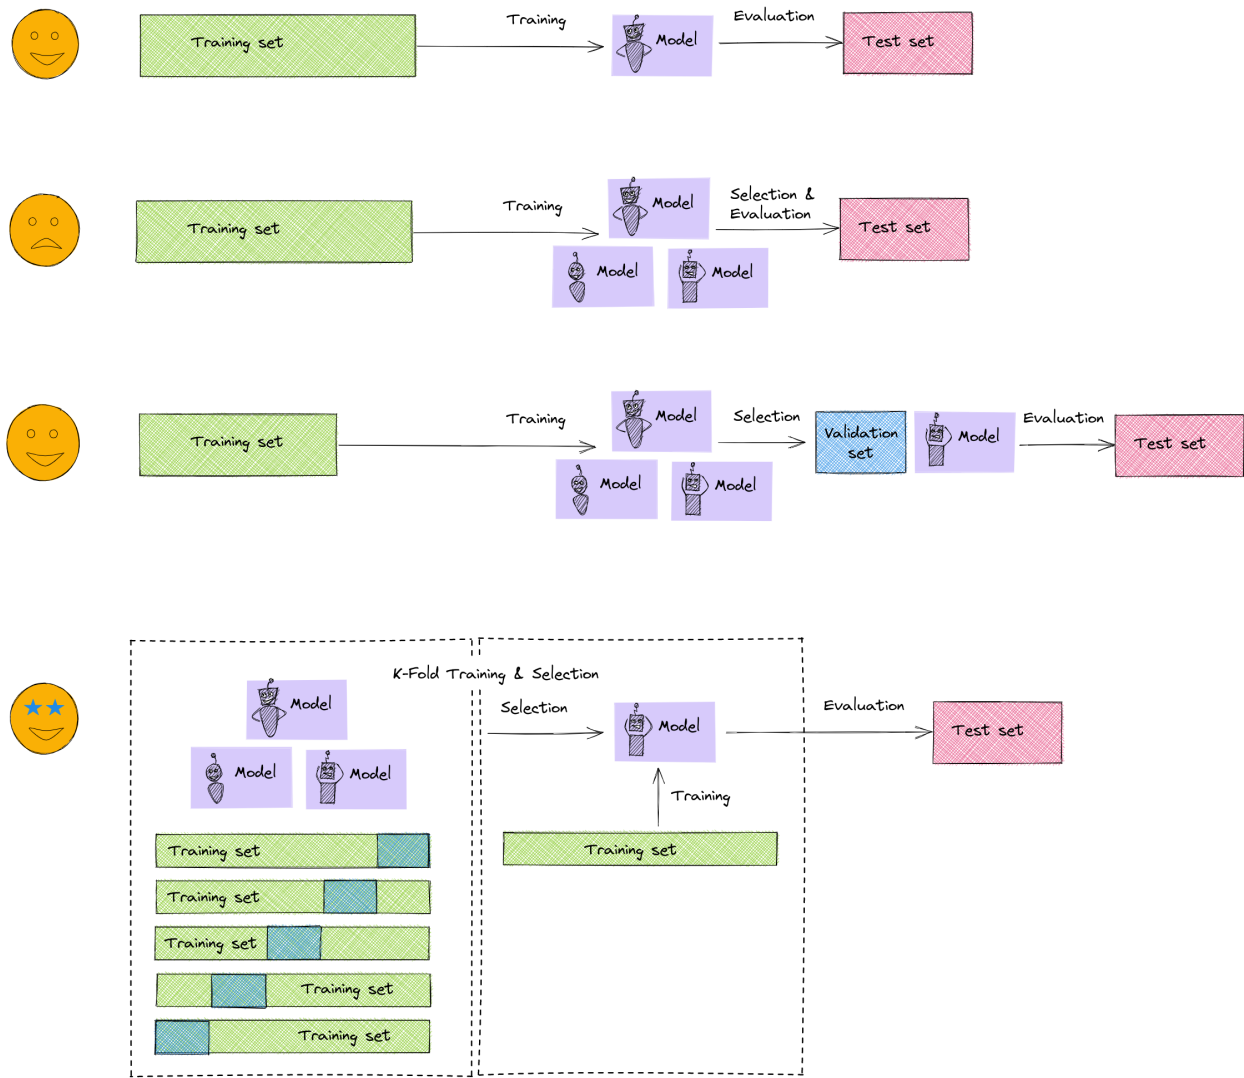


Figure 2: Summary of approaches. [Source](#).

## K-fold CV - how to split?

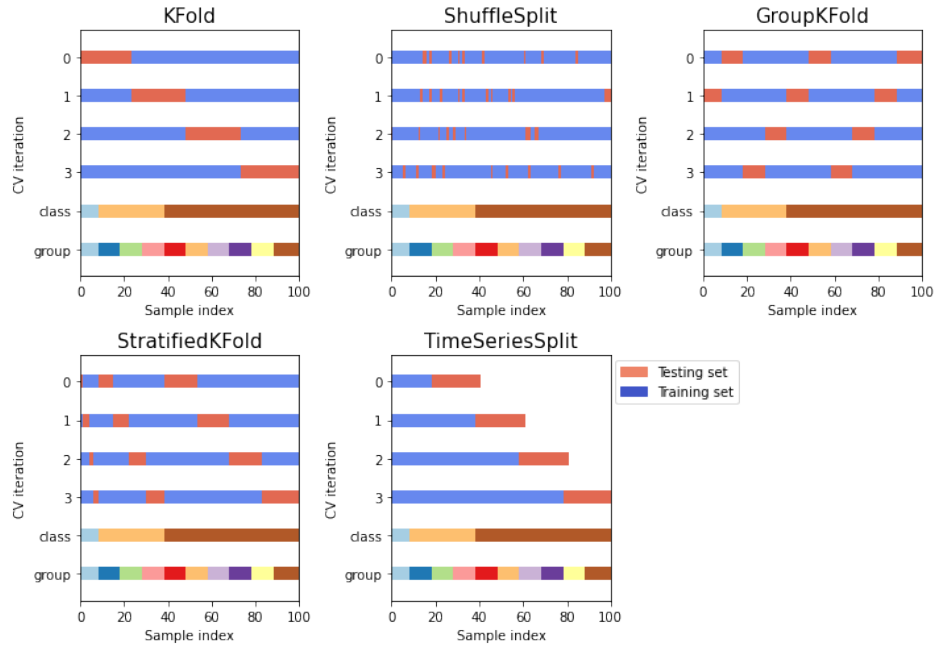


Figure 3: K-fold CV variations.

Selecting the right K-fold CV is very important for avoiding data leakage!

Refer to [the function documentation](#) for more examples.

### One standard error rule

- Model selection that minimizes mean error often results in too-complex model
- One standard error rule: use simplest model where mean error is within one SE of the minimum mean error

### One standard error rule - algorithm (1)

- Given data  $X, y$
- Compute score  $S_{p,i}$  for model  $p$  on fold  $i$  (of  $K$ )
- Compute average ( $\bar{S}_p$ ), standard deviation  $\sigma_p$ , and standard error of scores:

$$SE_p = \frac{\sigma_p}{\sqrt{K-1}}$$

### One standard error rule - algorithm (2)

“Best score” model selection:  $p^* = \operatorname{argmin}_p \bar{S}_p$

**One SE rule** for “lower is better” scoring metric: Compute target score:  $S_t = \bar{S}_{p^*} + SE_{p^*}$

then select simplest model with score lower than target:

$$p^{*,1SE} = \min\{p | \bar{S}_p \leq S_t\}$$

Note: this assumes you are using a “smaller is better” metric such as MSE. If you are using a “larger is better” metric, like R2, how would we change the algorithm?

### One standard error rule - algorithm (3)

“Best score” model selection:  $p^* = \operatorname{argmax}_p \bar{S}_p$

**One SE rule** for “higher is better” scoring metric: Compute target score:  $S_t = \bar{S}_{p^*} - SE_{p^*}$

then select simplest model with score higher than target:

$$p^{*,1SE} = \min\{p | \bar{S}_p \geq S_t\}$$