

K Nearest Neighbor

Fraida Fund

Contents

In this lecture	2
Parametric vs. non-parametric models	3
So far...	3
Flexible models	3
Parametric models	3
Non-parametric models	3
Nearest neighbor	3
1-NN	3
1NN - decision boundaries	3
K nearest neighbors	3
KNN for classification (1)	3
KNN for classification (Illustration)	4
KNN for classification (2)	4
KNN for classification (3)	4
KNN for regression	5
Model choices	5
What value of K?	5
What value of K? Illustration (1NN)	5
What value of K? Illustration (2NN)	7
What value of K? Illustration (3NN)	7
What value of K? Illustration (9NN)	7
What distance measure? (1)	7
What distance measure? (2)	7
Standardization (1)	7
Standardization (2)	8
What distance measure? (3)	8
How to combine labels into prediction?	8
Nearest neighbor in <code>sklearn</code>	9
The Curse of Dimensionality	9
Space grows exponentially with dimension	9
KNN in 1D	9
KNN in 2D	9
Density of samples decreases with dimensions	9
Density of samples decreases with dimensions - general	10
Density of samples decreases with dimensions - illustration	10
Solutions to the curse (1)	10
Solutions to the curse (2)	10
Bias and variance of KNN	11
True function	11
Expected loss	11
KNN output	11
Bias of KNN	11

Variance of KNN (1)	12
Variance of KNN (2)	12
MSE of KNN	12
Bias variance tradeoff	12
Bias variance tradeoff - selected problems	13
Summary of NN method	13
NN learning	13
NN prediction	13
The good and the bad (1)	13
The good and the bad (2)	13
The good and the bad (3)	13
Illustration	13

In this lecture

- Parametric vs. non-parametric models
- Nearest neighbor
- Model choices
- The Curse of Dimensionality
- Bias and variance of KNN

Parametric vs. non-parametric models

So far...

All of our models have looked like

$$\hat{y} = \beta_0 + \beta_1 x_1 + \dots + \beta_p x_p$$

Flexible models

A model is more flexible if it can produce more possibilities for $f(\mathbf{x})$.

- Some possibilities for $f(x_1, \dots, x_p) = \beta_0 + \beta_1 x_1$
- More possibilities for $f(x_1, \dots, x_p) = \beta_0 + \beta_1 x_1 + \beta_2 x_2$
- Even more possibilities for $f(x_1, \dots, x_p) = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \beta_3 x_3$

A way to get more flexible models is with a **non-parametric** approach.

Parametric models

- A particular model class is assumed (e.g. linear)
- Number of parameters fixed in advance

Non-parametric models

- Minimal assumptions about model class
- Model structure determined by data

Nearest neighbor

- A kind of non-parametric model.
- Basic idea: Find labeled samples that are “similar” to the new sample, and use their labels to make prediction for the new sample.

1-NN

- Given training data $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_N, y_N)$ and a new sample \mathbf{x}_0
- Find the sample in the training data $\mathbf{x}_{i'}$ with the least distance to \mathbf{x}_0 .

$$i' = \underset{i=1, \dots, N}{\operatorname{argmin}} d(\mathbf{x}_i, \mathbf{x}_0)$$

- Let $\hat{y}_0 = y_{i'}$

1NN - decision boundaries

K nearest neighbors

Instead of 1 closest sample, we find K closest samples.

KNN for classification (1)

Idea: Estimate conditional probability for class as fraction of points among neighbors with that class label.




Figure 1: 1NN decision boundaries - Nearest neighborhoods for each point of the training data set. Note that there will be zero error on training set!

KNN for classification (Illustration)

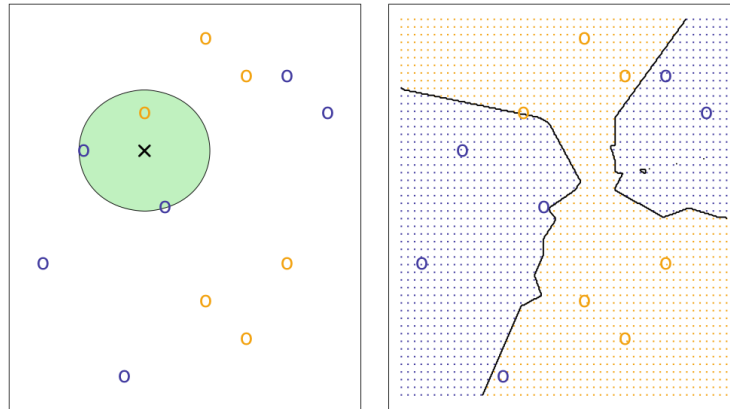


Figure 2: 3-NN classification. Image from ISLR.

KNN for classification (2)

Let N_0 be the set of K points in the training data that are closest to \mathbf{x}_0 . Then, for each class $m \in M$, we can estimate the conditional probability

$$P(y = m | \mathbf{x}_0) = \frac{1}{K} \sum_{(\mathbf{x}_i, y_i) \in N_0} I(y_i = m)$$

where $I(y_i = m)$ is an indicator variable that evaluates to 1 if for a given observation $(\mathbf{x}_i, y_i) \in N_0$ is a member of class m , and 0 otherwise.

KNN for classification (3)

- We can then select the class with the highest probability.

- Practically: select the most frequent class among the neighbors.

KNN for regression

Idea: Use the the combined label of the K nearest neighbors.

Let N_0 be the set of K points in the training data that are closest to \mathbf{x}_0 . Then,

$$\hat{y}_0 = \frac{1}{K} \sum_{(\mathbf{x}_i, y_i) \in N_0} y_i$$

Model choices

- What value of K ?
- What distance measure?
- How to combine K labels into prediction?

What value of K?

- K must be specified, can be selected by CV. (In general: larger K, less complex model.)
- Often cited “rule of thumb”: use $K = \sqrt{N}$
- Alternative: Radius-based neighbor learning, where fixed radius r is specified, can be selected by CV. (Number of neighbors depends on local density of points.)

What value of K? Illustration (1NN)




Figure 3: 1NN




Figure 4: 2NN




Figure 5: 3NN




Figure 6: 9NN

What value of K? Illustration (2NN)

What value of K? Illustration (3NN)

What value of K? Illustration (9NN)

What distance measure? (1)

- Euclidean (L2): $\sqrt{\sum_{i=1}^k (a_i - b_i)^2}$
- Manhattan (L1): $\sum_{i=1}^k |a_i - b_i|$
- Minkowski (generalization of both): $\left(\sum_{i=1}^k (|a_i - b_i|)^q \right)^{\frac{1}{q}}$

(L2 distance prefers many medium disagreements to one big one.)

What distance measure? (2)

- Feature weight depends on scale
- KNN implicitly weights all features equally

Standardization (1)




Figure 7: Without standardization, via <https://stats.stackexchange.com/a/287439/>

Standardization (2)




Figure 8: With standardization, via <https://stats.stackexchange.com/a/287439/>

What distance measure? (3)

- Alternative to equal weighted features: assign feature weights

$$d(\mathbf{a}, \mathbf{b}) = \left(\sum_{i=1}^k (w_i |a_i - b_i|)^q \right)^{\frac{1}{q}}$$

How to combine labels into prediction?

- Basic voting: use mode of neighbors for classification, mean or median for regression.
- Distance-weighted: weight of vote inversely proportional to distance from the query point. ("More similar" training points count more.)

Nearest neighbor in sklearn

- `KNeighborsClassifier`
- `KNeighborsRegressor`

The Curse of Dimensionality

Space grows exponentially with dimension




Figure 9: Number of cells grows exponentially with dimension. From Bishop PRML, Fig. 1-21

KNN in 1D

- Consider a dataset $(x_1, y_1), \dots, (x_N, y_N)$ with $N=100$
- x takes on values in the range $[0,1]$ with uniform distribution
- On average, one data point is located every $1/100$ units along 1D feature axis.
- To find 3NN, would expect to cover $3/100$ of the feature axis.

KNN in 2D

- Now consider a dataset with two features: $([x_{1,1}, x_{1,2}], y_1), \dots, ([x_{N,1}, x_{N,2}], y_N)$ with $N=100$
- Each feature takes on values in the range $[0,1]$ with uniform distribution
- To find 3NN, would expect to cover $0.03^{\frac{1}{2}}$ of the unit rectangle.

Density of samples decreases with dimensions

To get 3NN,

- need to cover 3% of space in 1D
- need to cover 17% of space in 2D
- need to cover 70% of space in 10D. At this point, the nearest neighbors are not much closer than the rest of the dataset.

Density of samples decreases with dimensions - general

The length of the smallest hyper-cube that contains all K-nearest neighbors of a test point:

$$\left(\frac{K}{N}\right)^{\frac{1}{d}}$$

for N samples with dimensionality d .

Density of samples decreases with dimensions - illustration




Figure 10: Image source: https://www.cs.cornell.edu/courses/cs4780/2018fa/lectures/lecturenote02_kNN.html

Solutions to the curse (1)

Add training data?

$$\left(\frac{K}{N}\right)^{\frac{1}{d}}$$

As number of dimensions increases linearly, number of training samples must increase exponentially to counter the “curse”.

Solutions to the curse (2)

Reduce d ?

- Feature selection (and in previous lab)
- Dimensionality reduction: a type of unsupervised learning that *transforms* high-d data into lower-d data.

Bias and variance of KNN

True function

Suppose data has true relation

$$y = f_0(\mathbf{x}) + \epsilon, \quad \epsilon \sim N(0, \sigma_\epsilon^2)$$

and our model predicts $\hat{y} = f(\mathbf{x})$.

Expected loss

We will use least square loss function, so that the expected error at a given test point \mathbf{x}_{test} is:

$$\begin{aligned} MSE_y(\mathbf{x}_{test}) &:= E[y - \hat{y}]^2 \\ &= E[f_0(\mathbf{x}_{test}) + \epsilon_0 - f(\mathbf{x}_{test})]^2 \\ &= E[f_0(\mathbf{x}_{test}) - f(\mathbf{x}_{test})]^2 + \sigma_\epsilon^2 \end{aligned}$$

which we know can be further decomposed into squared bias, variance, and irreducible error.

KNN output

The output of the KNN algorithm at a test point is

$$f(\mathbf{x}_{test}) = \frac{1}{K} \sum_{\ell \in K_x} f_0(\mathbf{x}_\ell) + \epsilon_\ell$$

where K_x is the set of K nearest neighbors of \mathbf{x}_{test} . (We assume that these neighbors are fixed.)

Bias of KNN

When we take the expectation of bias over all test samples, the ϵ values disappear because it has zero mean, so squared bias becomes

$$\begin{aligned} Bias^2 &= (f_0(\mathbf{x}_{test}) - E[f(\mathbf{x}_{test})])^2 \\ &= \left(f_0(\mathbf{x}_{test}) - E \left(\frac{1}{K} \sum_{\ell \in K_x} f_0(\mathbf{x}_\ell) + \epsilon_\ell \right) \right)^2 \\ &= \left(f_0(\mathbf{x}_{test}) - \frac{1}{K} \sum_{\ell \in K_x} f_0(\mathbf{x}_\ell) \right)^2 \end{aligned}$$

(We also rely on the assumption that neighbors are fixed.)

Variance of KNN (1)

$$\begin{aligned} Var(\hat{y}) &= Var\left(\frac{1}{K} \sum_{\ell \in K_x} f_0(\mathbf{x}_\ell) + \epsilon_\ell\right) \\ &= \frac{1}{K^2} \sum_{\ell \in K_x} Var(f_0(x_\ell) + \epsilon_\ell) \\ &= \frac{1}{K^2} \sum_{\ell \in K_x} Var(f_0(x_\ell)) + Var(\epsilon_\ell) \end{aligned}$$

Variance of KNN (2)

$$\begin{aligned} &= \frac{1}{K^2} \sum_{\ell \in K_x} Var(\epsilon_\ell) \\ &= \frac{1}{K^2} K \sigma_\epsilon^2 \\ &= \frac{\sigma_\epsilon^2}{K} \end{aligned}$$

where $Var(f_0(x_\ell)) = 0$ only if we assume that the neighbors are fixed.

MSE of KNN

Then the MSE of KNN is

$$MSE(\mathbf{x}_{test}) = \left(f_0(\mathbf{x}_{test}) - \frac{1}{K} \sum_{\ell \in K_x} f_0(\mathbf{x}_\ell) \right)^2 + \frac{\sigma_\epsilon^2}{K} + \sigma_\epsilon^2$$

where K_x is the set of K nearest neighbors of \mathbf{x}_{test} .

Bias variance tradeoff

- Variance decreases with K
- Bias likely to increase with K, if function $f_0(\cdot)$ is smooth. (Few closest neighbors to test point will have similar values, so average will be close to $f_0(\mathbf{x})$; as K increases, neighbors are further away, and average of neighbors moves away from $f_0(\mathbf{x})$.)

Bias variance tradeoff - selected problems

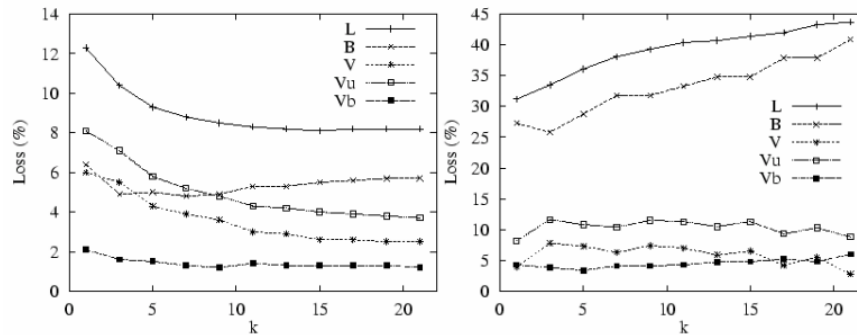


Figure 11: Via Domingos 2000.

Summary of NN method

NN learning

Learning:

- Store training data
- Don't do anything else until you have a new point to classify

NN prediction

Prediction:

- Find nearest neighbors using distance metric
- Classification: predict most frequently occurring class among nearest neighbors
- Regression: predict mean value of nearest neighbors

The good and the bad (1)

Good:

- Good interpretability
- Fast "learning" (*memory-based*)
- Works well in low dimensions for complex decision surfaces

The good and the bad (2)

Neutral:

- Assumes similar inputs have similar outputs

The good and the bad (3)

Bad:

- Slow prediction (especially with large N)
- Curse of dimensionality

Illustration




Figure 12: 1NN, 3NN, 9NN comparison on three types of data, with accuracy on test set shown in the corner.