Logistic Regression for Classification

Fraida Fund

Contents

In this lecture

- Linear classifiers
- Logistic regression
- Fitting logistic regression
- Naive Bayes classifier

Classification

Suppose we have a series of data points $\{(\mathbf{x_1}, y_1), (\mathbf{x_2}, y_2), \dots, (\mathbf{x_n}, y_n)\}$ and there is some (unknown) relationship between $\mathbf{x_i}$ and y_i.

- **Classification**: The output variable y is constrained to be $\in 1, 2, \cdots, K$

- **Binary classification**: The output variable y is constrained to be $\in 0, 1$

Linear classifiers

Binary classification with linear decision boundary

- Plot training data points
- Draw a line (**decision boundary**) separating 0 class and 1 class
- If a new data point is in the **decision region** corresponding to class 0, then $\hat{y} = 0$.
- If it is in the decision region corresponding to class 1, then $\hat{y} = 1$.
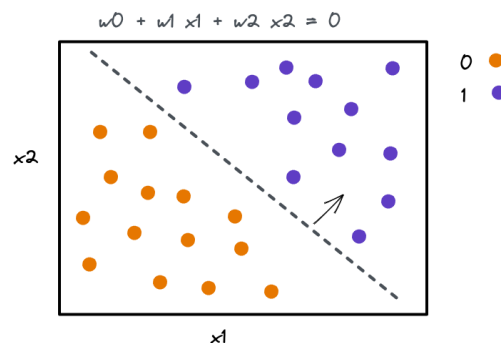
Figure 1: Binary classification problem with linear decision boundary.

Linear classification rule

- Given a **weight vector**: $\mathbf{w} = (w_0, \cdots, w_d)$
- Compute linear combination $z = w_0 + \sum_{j=1}^{d} w_d x_d$
- Predict class:

$$\hat{y} = \begin{cases} 1, z > 0 \\ 0, z \leq 0 \end{cases}$$

Multi-class classification: illustration
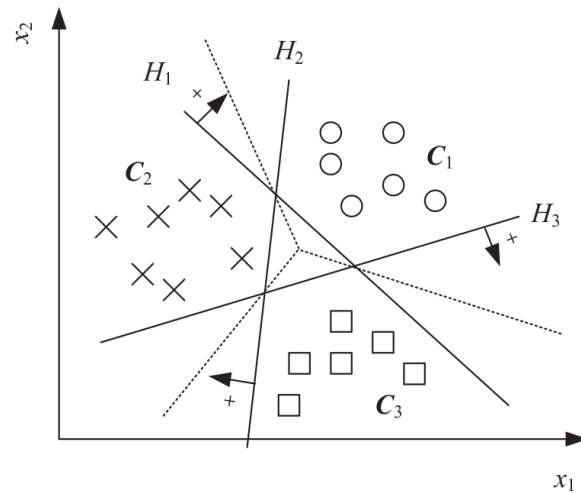
Figure 2: Each hyperplane H_i separates the examples of C_i from the examples of all other classes.

Linear separability

Given training data

$$(\mathbf{x}_i, y_i), i = 1, \cdots, N$$

The problem is **perfectly linearly separable** if there exists a **separating hyperplane** H_i such that all $\mathbf{x} \in C_i$ lie on its positive side, and all $\mathbf{x} \in C_j, j \neq i$ lie on its negative side.

Non-uniqueness of separating hyperplane

When a separating hyperplane exists, it is not unique (there are in fact infinitely many such hyperplanes.)
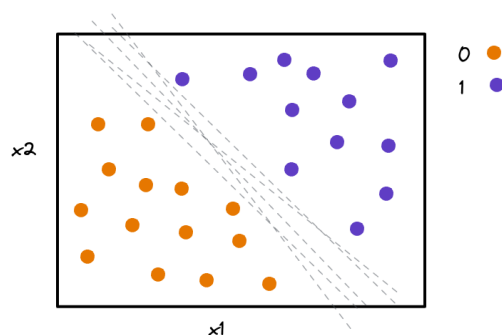
Figure 3: Several separating hyperplanes.

Non-existence of perfectly separating hyperplane

Many datasets *not* linearly separable - some points will be misclassified by *any* possible hyperplane.
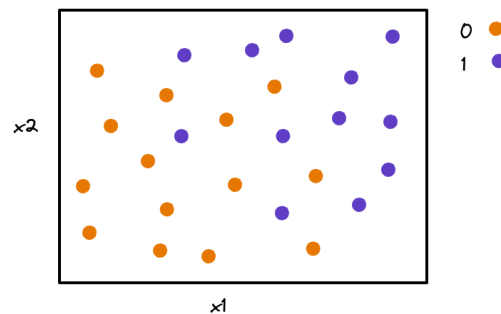
Figure 4: This data is not separable.

Choosing a hyperplane

Which hyperplane to choose?

We will try to find the hyperplane that minimizes loss according to some **loss function**.

Will revisit several times this semester.

Logistic regression

Probabilistic model for binary classification

Instead of looking for a model f so that

$$y_i \approx f(x_i)$$

we will look for an f so that

$$P(y_i = 1|x_i) = f(x_i), P(y_i = 0|x_i) = 1 - f(x_i)$$

We need a function that takes a real value and maps it to range $[0, 1]$. What function should we use?

Logistic/sigmoid function
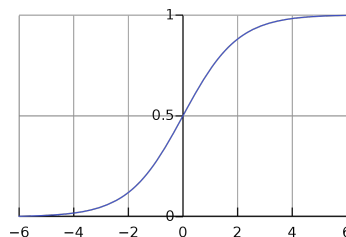
Figure 5: $\sigma(z) = \frac{1}{1+e^{-z}}$ is a classic "S"-shaped function.

Note the intuitive relationship behind this function's output and the distance from the linear separator (the argument that is input to the function).

5
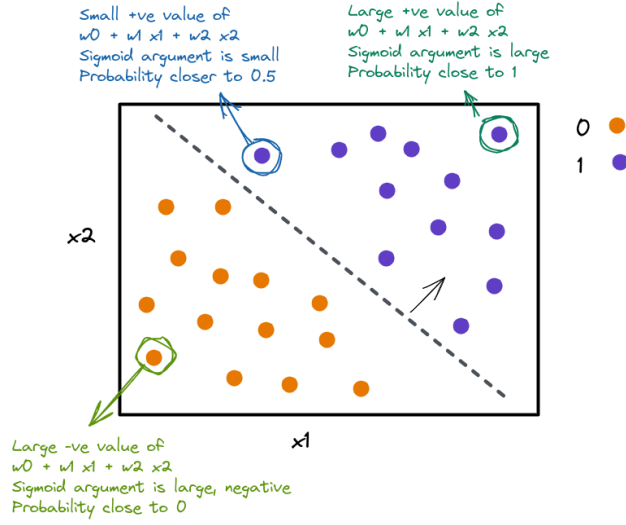
Figure 6: Output is close to 0 or 1 if the argument to the σ has large magnitude (point is far from separating hyperplane, but closer to 0.5 if the argument is small (point is near separating hyperplane).

Logistic function for binary classification

Let $z = w_0 + \sum_{j=1}^{d} w_d x_d$, then

$$P(y = 1|\mathbf{x}) = \frac{1}{1 + e^{-z}}, \quad P(y = 0|\mathbf{x}) = \frac{e^{-z}}{1 + e^{-z}}$$

(note: $P(y = 1) + P(y = 0) = 1$)

Logistic function with threshold

Choose a threshold t, then

$$\hat{y} = \begin{cases} 1, & P(y = 1|\mathbf{x}) \geq t \\ 0, & P(y = 1|\mathbf{x}) < t \end{cases}$$

Logistic model as a "soft" classifier
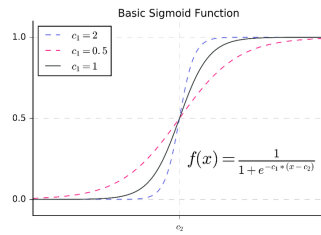
Figure 7: Plot of $P(y = 1|x) = \frac{1}{1+e^{-z}}$, $z = w_1 x$. As $w_1 \to \infty$ the logistic model becomes a "hard" rule.

Logistic classifier properties (1)

- Class probabilities depend on distance from separating hyperplane
- Points far from separating hyperplane have probability ≈ 0 or ≈ 1
- When $||\mathbf{w}||$ is larger, class probabilities go towards extremes (0,1) more quickly

Logistic classifier properties (2)

- Unlike linear regression, weights do *not* correspond to change in output associated with one-unit change in input.
- Sign of weight *does* tell us about relationship between a given feature and target variable.

Logistic regression - illustration
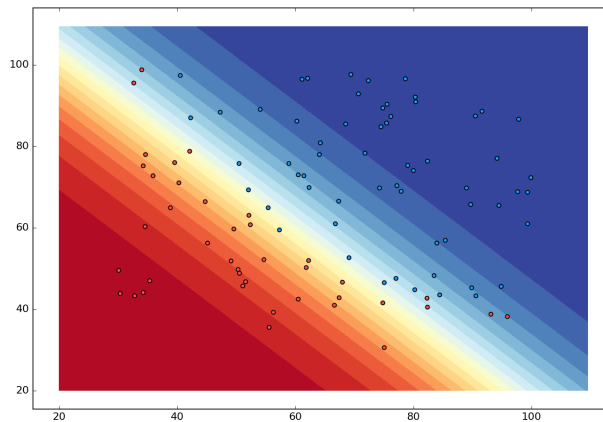
Figure 8: Logistic regression, illustrated with contour plot.

Multi-class logistic regression

Suppose $y \in 1, \ldots, K$. We use:

- $\mathbf{W} \in R^{K \times d}$ (parameter matrix)
- $\mathbf{z} = \mathbf{W}\mathbf{x}$ (K linear functions)

Assume we have stacked a 1s column so that the intercept is rolled into the parameter matrix.

Softmax function

$$g_k(\mathbf{z}) = \frac{e^{z_k}}{\sum_{\ell=1}^{K} e^{z_\ell}}$$

- Takes as input a vector of K numbers
- Outputs K probabilities proportional to the exponentials of the input numbers.

Softmax function as a PMF

Acts like a probability mass function:

- $g_k(\mathbf{z}) \in [0, 1]$ for each k
- $\sum_{k=1}^{K} g_k(\mathbf{z}) = 1$
- larger input corresponds to larger "probability"

Softmax function for multi-class logistic regression (1)

Class probabilities are given by

$$P(y = k|\mathbf{x}) = \frac{e^{z_k}}{\sum_{\ell=1}^{K} e^{z_\ell}}$$

Softmax function for multi-class logistic regression (2)

When $z_k \gg z_\ell$ for all $\ell \neq k$:

- $g_k(\mathbf{z}) \approx 1$
- $g_\ell(\mathbf{z}) \approx 0$ for all $\ell \neq k$

Assign highest probability to class k when z_k is largest.

Fitting logistic regression model

We know that to fit weights, we need

- a loss function,
- and a training algorithm to find the weights that minimize the loss function.

Learning logistic model parameters

Weights $\mathbf{W}$ are the unknown **model parameters**:

$$\mathbf{z} = \mathbf{W}\mathbf{x}, \mathbf{W} \in R^{K \times d}$$

$$P(y = k|\mathbf{x}) = g_k(\mathbf{z}) = g_k(\mathbf{W}\mathbf{x})$$

Given training data $(\mathbf{x}_i, y_i), i = 1, \dots, n$, we must learn $\mathbf{W}$.

Note that if the data is linearly separable, there will be more than one $\mathbf{W}$ that perfectly classifies the training data! We will choose the *maximum likelihood* one.

Maximum likelihood estimation (1)

Let $P(\mathbf{y}|\mathbf{X}, \mathbf{W})$ be the probability of observing class labels $\mathbf{y} = (y_1, \dots, y_n)^T$

given inputs $\mathbf{X} = (\mathbf{x}_1, \dots, \mathbf{x}_n)^T$ and weights $\mathbf{W}$.

The **maximum likelihood estimate** is

$$\hat{\mathbf{W}} = \underset{W}{\operatorname{argmax}} P(\mathbf{y}|\mathbf{X}, \mathbf{W})$$

It is the estimate of parameters for which these observations are most likely.

Maximum likelihood estimation (2)

Assume outputs y_i are independent of one another,

$$P(\mathbf{y}|\mathbf{X}, \mathbf{W}) = \prod_{i=1}^{n} P(y_i|\mathbf{x_i}, \mathbf{W})$$

Note: $P(y_i|\mathbf{x_i}, \mathbf{W})$ is equal to

- $y_i P(y_i = 1|\mathbf{x_i}, \mathbf{W})$ when $y_i = 1$
- and $(1 - y_i)P(y_i = 0|\mathbf{x_i}, \mathbf{W})$ when $y_i = 0$.

and since only one term will be non-zero for any given y_i, $P(y_i|\mathbf{x_i}, \mathbf{W})$ is equal to the sum of those:

$$y_i P(y_i = 1|\mathbf{x_i}, \mathbf{W}) + (1 - y_i)P(y_i = 0|\mathbf{x_i}, \mathbf{W})$$

This expression is familiar as the PMF of a Bernoulli random variable.

We take the log of both sides, because then the product turns into a sum...

Maximum likelihood estimation (3)

Define the **negative log likelihood**:

$$L(\mathbf{W}) = -\ln P(\mathbf{y}|\mathbf{X}, \mathbf{W})$$
$$= -\sum_{i=1}^{n} \ln P(y_i|\mathbf{x_i}, \mathbf{W})$$

Note that maximizing the likelihood is the same as minimizing the negative log likelihood.

Maximum likelihood estimation (4)

Now we can re-write max likelihood estimator with a loss function to minimize:

$$\hat{\mathbf{W}} = \underset{W}{\operatorname{argmax}}\, P(\mathbf{y}|\mathbf{X}, \mathbf{W}) = \underset{W}{\operatorname{argmin}}\, L(\mathbf{W})$$

At this point, we know we need to find

$$\underset{W}{\operatorname{argmin}}\left(-\sum_{i=1}^{n} y_i \ln P(y_i = 1|\mathbf{x_i}, \mathbf{W}) + (1 - y_i) \ln P(y_i = 0|\mathbf{x_i}, \mathbf{W})\right)$$

The next step will be to plug in our sigmoid function, $P(y_i = 1|\mathbf{x_i}, \mathbf{W}) = \sigma(z_i)$ where $z_i = \mathbf{W}\mathbf{x_i}$.

Binary cross-entropy loss (1)

For binary classification with class labels $0, 1$:

$$\begin{aligned}
\ln P(y_i|\mathbf{x_i}, \mathbf{w}) &= y_i \ln P(y_i = 1|\mathbf{x_i}, \mathbf{w}) + (1 - y_i) \ln P(y_i = 0|\mathbf{x_i}, \mathbf{w}) \\
&= y_i \ln \sigma(z_i) + (1 - y_i) \ln(1 - \sigma(z_i)) \\
&= y_i(\ln \sigma(z_i) - \ln \sigma(-z_i)) + \ln \sigma(-z_i) \\
&= y_i \ln \frac{\sigma(z_i)}{\sigma(-z_i)} + \ln \sigma(-z_i) \\
&= y_i \ln \frac{1 + e^{z_i}}{1 + e^{-z_i}} + \ln \sigma(-z_i) \\
&= y_i \ln \frac{e^{z_i}(e^{-z_i} + 1)}{1 + e^{-z_i}} + \ln \sigma(-z_i) \\
&= y_i z_i - \ln(1 + e^{z_i})
\end{aligned} \tag{1}$$

Notes: $\sigma(-z) = 1 - \sigma(z)$

Binary cross-entropy loss (2)

Binary cross-entropy loss function (negative log likelihood) for $[0, 1]$ class labels:

$$-\sum_{i=1}^{n} \ln P(y_i|\mathbf{x_i}, \mathbf{W}) = \sum_{i=1}^{n} \ln(1 + e^{z_i}) - y_i z_i$$

Cross-entropy loss for multi-class classification (1)

Define "one-hot" vector - for a sample from class k, all entries in the vector are 0 except for the kth entry which is 1:

$$r_{ik} = \begin{cases} 1 & y_i = k \\ 0 & y_i \neq k \end{cases}$$

$$i = 1, \ldots, n, \quad k = 1, \ldots, K$$

For example: if the class labels are $[0, 1, 2, 3, 4]$, for a sample for which $y_i = 3$, $r_{ik} = [0, 0, 0, 1, 0]$.

Cross-entropy loss for multi-class classification (2)

Then,

$$\ln P(y_i | \mathbf{x_i}, \mathbf{W}) = \sum_{k=1}^{K} r_{ik} \ln P(y_i = k | \mathbf{x_i}, \mathbf{W})$$

Cross-entropy loss function is

$$\sum_{i=1}^{n} \left[\ln \left(\sum_k e^{z_{ik}} \right) - \sum_k z_{ik} r_{ik} \right]$$

Minimizing cross-entropy loss

To minimize, we would take the partial derivative:

$$\frac{\partial L(W)}{\partial W_{kj}} = 0$$

for all W_{kj}

But, there is no closed-form expression - can only estimate weights via numerical optimization (e.g. gradient descent)

Non-linear decision boundaries

- Logistic regression learns linear boundary
- What if the "natural" decision boundary is non-linear?

Can use basis functions to map problem to transformed feature space (if "natural" decision boundary is non-linear)

Bias, variance

- Variance increases with d and decreases with n
- Can add a regularization penalty to loss function

"Recipe" for logistic regression (binary classifier)

- Choose a **model**:

$$P(y = 1|x, w) = \sigma \left(w_0 + \sum_{i=1}^{d} w_d x_d \right)$$

$$\hat{y} = \begin{cases} 1, & P(y = 1|\mathbf{x}) \geq t \\ 0, & P(y = 1|\mathbf{x}) < t \end{cases}$$

- Get **data** - for supervised learning, we need **labeled** examples: $(x_i, y_i), i = 1, 2, \cdots, n$
- Choose a **loss function** that will measure how well model fits data: binary cross-entropy

$$\sum_{i=1}^{n} \ln(1 + e^{z_i}) - y_i z_i$$

- Find model **parameters** that minimize loss: use numerical optimization to find weight vector w
- Use model to **predict** $\hat{y}$ for new, unlabeled samples.

"Recipe" for logistic regression (multi-class classifier)

- Choose a **model**: find probability of belonging to each class, then choose the class for which the probability is highest.

$$P(y = k|\mathbf{x}) = \frac{e^{z_k}}{\sum_{\ell=1}^{K} e^{z_\ell}} \text{ where } \mathbf{z} = \mathbf{W}\mathbf{x}$$

- Get **data** - for supervised learning, we need **labeled** examples: $(x_i, y_i), i = 1, 2, \cdots, n$
- Choose a **loss function** that will measure how well model fits data: categorical cross-entropy

$$\sum_{i=1}^{n} \left[\ln \left(\sum_k e^{z_{ik}} \right) - \sum_k z_{ik} r_{ik} \right] \text{ where}$$

$$r_{ik} = \begin{cases} 1 & y_i = k \\ 0 & y_i \neq k \end{cases}$$

- Find model **parameters** that minimize loss: use numerical optimization to find weight vector w
- Use model to **predict** $\hat{y}$ for new, unlabeled samples.

Naive Bayes classifier

A quick look at a different type of model!

Probabilistic models (1)

For logistic regression, minimizing the cross-entropy loss finds the parameters for which

$$P(\mathbf{y}|\mathbf{X}, \mathbf{W})$$

is maximized.

Probabilistic models (2)

For linear regression, assuming normally distributed stochastic error, minimizing the **squared error** loss finds the parameters for which

$$P(\mathbf{y}|\mathbf{X}, \mathbf{w})$$

is maximized.

Surprise! We've been doing maximum likelihood estimation all along.

Probabilistic models (3)

ML models that try to

- get a good fit for $P(y|X)$: **discriminative** models.
- fit $P(X, y)$ or $P(X|y)P(y)$: **generative** models.

Linear regression and logistic regression are both considered discriminative models; they say "given that we have this data, what's the most likely label?" (e.g. learning a mapping from an input to a target variable).

Generative models try to learn "what does data for each class look like" and then apply Bayes rule.

Bayes rule

For a sample $\mathbf{x}_i$, y_k is label of class k:

$$P(y_k|\mathbf{x}_i) = \frac{P(\mathbf{x}_i|y_k)P(y_k)}{P(\mathbf{x}_i)}$$

- $P(y_k|\mathbf{x}_i)$: posterior probability. "What is the probability that this sample belongs to class k, given its observed feature values are $\mathbf{x}_i$?"
- $P(\mathbf{x}_i|y_k)$: conditional probability: "What is the probability of observing the feature values $\mathbf{x}_i$ in a sample, given that the sample belongs to class k?"
- $P(y_k)$: prior probability
- $P(\mathbf{x}_i)$: evidence

Class conditional probability (1)

"Naive" assumption conditional independence of features:

$$P(\mathbf{x}_i|y_k) = P(x_{i,1}|y_k)P(x_{i,2}|y_k)\,...\,P(x_{i,d}|y_k)$$

$$= \prod_{j=1}^{d} P(x_{i,j}|y_k)$$

This is called "naive" because this assumption is probably not true in most realistic situations.
(But the classifier may still work OK!)
Also assumes samples are i.i.d.

Class conditional probability (2)

Example: for binary/categorical features, we could compute

$$\hat{P}(x_{i,j}|y_k) = \frac{N_{x_{i,j},y_k}}{N_{y_k}}$$

- $N_{x_{i,j},y_k}$ is the number of samples belonging to class k that have feature j.
- N_{y_k} is the total number of samples belonging to class k.

Example: for cat photo classifier,

$$\hat{P}(\mathbf{x}_i = [\text{has tail, has pointy ears, has fur, purrs when petted, likes to eat fish}]|y = \text{cat})$$

$$\rightarrow P(\frac{N_{\text{tail, cat}}}{N_{\text{cat}}})P(\frac{N_{\text{pointy ears, cat}}}{N_{\text{cat}}})P(\frac{N_{\text{fur, cat}}}{N_{\text{cat}}})P(\frac{N_{\text{purrs, cat}}}{N_{\text{cat}}})P(\frac{N_{\text{eats fish, cat}}}{N_{\text{cat}}})$$

$$\rightarrow \frac{20}{20}\frac{18}{20}\frac{17}{20}\frac{5}{20}\frac{15}{20}$$

Prior probability

Can estimate prior probability as

$$\hat{P}(y_k) = \frac{N_{y_k}}{N}$$

Prior probabilities: probability of encountering a particular class k.

Example: $\frac{20}{1500}$ photos are cats.

Evidence

We don't actually need $P(\mathbf{x}_i)$ to make decisions, since it is the same for every class.

Naive bayes decision boundary

Uniform prior (both classes equally likely)
Decision boundary where $P(x \mid 1) = P(x \mid 0)$

Purple class more likely: shifts distribution "up"
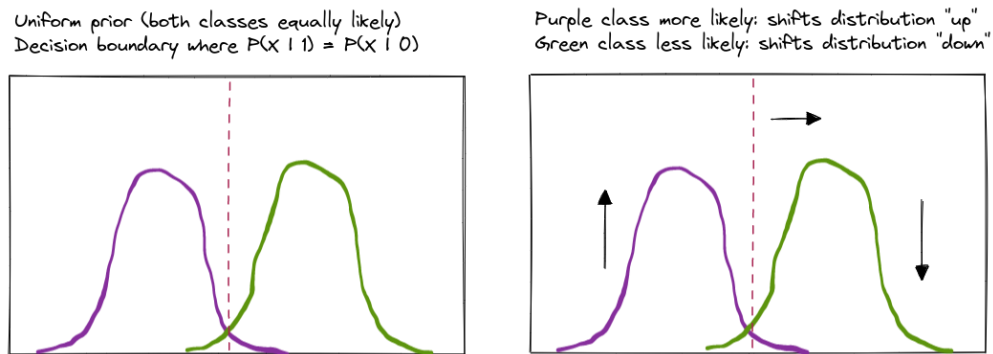Green class less likely: shifts distribution "down"

Figure 9: Naive bayes decision boundary.

Why generative model?

The generative model solves a more general problem than the discriminative model!

But, only the generative model can be used to **generate** new samples similar to the training data.

Example: "generate a new sample that is probably a cat."