

# Support vector machines

Fraida Fund

## Contents

In this lecture . . . . .	2
Recap . . . . .	3
Classifying data that is not linearly separable . . . . .	3
Maximal margin classifier . . . . .	3
Binary classification problem . . . . .	3
Linear separability . . . . .	3
Separating hyperplane (1) . . . . .	3
Separating hyperplane (2) . . . . .	3
Using the hyperplane to classify . . . . .	3
Non-uniqueness . . . . .	4
Which separating hyperplane is best? . . . . .	4
Margin . . . . .	4
Maximal margin classifier . . . . .	4
Maximal margin classifier - illustration . . . . .	4
Support vectors . . . . .	5
Constructing the maximal margin classifier (1) . . . . .	5
Constructing the maximal margin classifier (2) . . . . .	5
Constructing the maximal margin classifier (3) . . . . .	5
Constructing the maximal margin classifier (4) . . . . .	5
Problems with MM classifier (1) . . . . .	6
Problems with MM classifier (2) . . . . .	7
Support vector classifier . . . . .	7
Basic idea . . . . .	7
Constructing the support vector classifier . . . . .	7
Constructing the support vector classifier (3) . . . . .	7
Constructing the support vector classifier (4) . . . . .	8
Illustration of effect of $C$ . . . . .	8
Support vector . . . . .	8
$C$ controls bias-variance tradeoff . . . . .	8
Important terminology note . . . . .	8
Constrained vs. Lagrange forms . . . . .	9
Loss + penalty expression (1) . . . . .	9
Loss + penalty representation (2) . . . . .	9
Loss + penalty representation (3) . . . . .	9
Compared to logistic regression . . . . .	9
Hinge loss vs. logistic regression . . . . .	9
Maximizing the margin . . . . .	10
Optimization review . . . . .	10
Constrained optimization . . . . .	10
Definition of Lagrangian . . . . .	10
Dual problem (with extra details not shown in class) . . . . .	10
KKT conditions (1) . . . . .	11

KKT conditions (2)	11
Active vs. inactive constraints	11
Active vs. inactive constraints (illustration)	11
Comment on notation	12
Support vector classifier as constrained optimization (1)	12
Support vector classifier as constrained optimization (extra details)	12
Support vector classifier as constrained optimization (2)	13
Support vector classifier as constrained optimization (3)	13
Support vector classifier as constrained optimization (4)	13
Why solve dual problem?	13
Correlation interpretation (1)	14
Correlation interpretation (2)	14
Support vector machines	14
Extension to non-linear decision boundary	14
SVM in transformed form (1)	14
SVM in transformed form (2)	14
Kernel trick	15
Kernel trick example	15
Kernel trick example - direct computation	15
Kernel trick example - computation using kernel	15
Kernel intuition	15
Linear kernel	16
Polynomial kernel	16
Radial basis function kernel	17
Infinite-dimensional feature space	17
Infinite-dimensional feature space (extra steps not shown in class)	17
Infinite-dimensional feature space (2)	18
Extension to regression	18
SVR illustration	18
Summary: SVM	18
Key expression	18
Key ideas	19

## In this lecture

- Maximal margin classifier
- Support vector classifier
- Solving constrained optimization to find coefficients
- Support vector machine with non-linear kernel

## Recap

### Classifying data that is not linearly separable

- Decision tree - complex decision boundary, fast prediction, often works best as part of ensemble
- KNN - complex decision boundary, slow prediction
- Logistic regression - only if you use basis function  $\phi()$  to transform data before applying model

## Maximal margin classifier

### Binary classification problem

- $N$  training samples  $\mathbf{x}_1, \dots, \mathbf{x}_N \in \mathbb{R}^p$
- Class labels  $y_1, \dots, y_N \in \{-1, 1\}$

### Linear separability

The problem is **perfectly linearly separable** if there exists a **separating hyperplane**  $H_i$  such that

- all  $\mathbf{x} \in C_i$  lie on its positive side, and
- all  $\mathbf{x} \in C_j, j \neq i$  lie on its negative side.

### Separating hyperplane (1)

The separating hyperplane has the property that for all  $i = 1, \dots, N$ ,

$$\beta_0 + \sum_{j=1}^p \beta_j x_{ij} > 0 \text{ if } y_i = 1$$

$$\beta_0 + \sum_{j=1}^p \beta_j x_{ij} < 0 \text{ if } y_i = -1$$

### Separating hyperplane (2)

Equivalently:

$$y_i \left( \beta_0 + \sum_{j=1}^p \beta_j x_{ij} \right) > 0 \quad (1)$$

### Using the hyperplane to classify

Then, we can classify a new sample  $\mathbf{x}$  using the sign of

$$z = \beta_0 + \sum_{j=1}^p \beta_j x_{ij}$$

and we can use the magnitude of  $z$  to determine how confident we are about our classification. (Larger  $z$  = farther from hyperplane = more confident about classification.)

## Non-uniqueness

If a separating hyperplane exists, there will be an infinite number of separating hyperplanes.

## Which separating hyperplane is best?

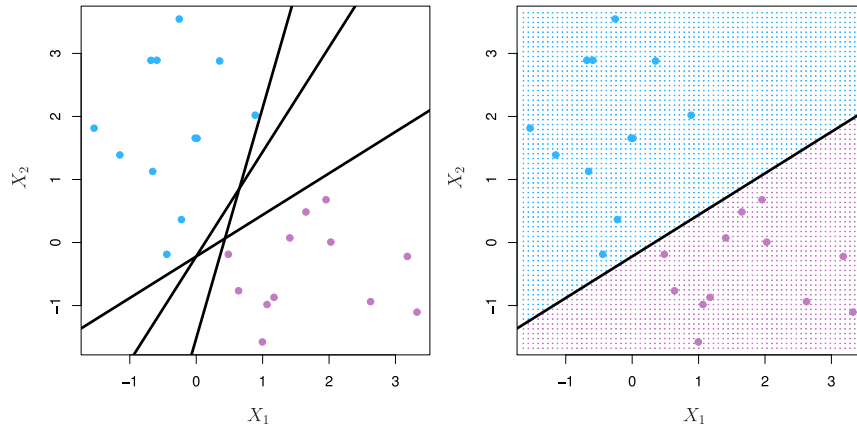


Figure 1: Fig. 9.2 from ISLR.

## Margin

- Compute distance from each training sample to the separating hyperplane.
- Smallest distance among all samples is called the **margin**.

## Maximal margin classifier

- For classifier to be more robust to noise, we should maximize the margin.
- Find the widest “slab” we can fit between the two classes.
- Choose the midline of this “slab” as the decision boundary.

## Maximal margin classifier - illustration




Figure 2: Fig. 9.3 from ISLR.

## Support vectors

- Points that lie on the border of maximal margin hyperplane are **support vectors**
- They “support” the maximal margin hyperplane: if these points move, then the maximal margin hyperplane moves
- Maximal margin hyperplane is not affected by movement of any other point, as long as it doesn't cross borders!

## Constructing the maximal margin classifier (1)

$$\underset{\beta, \gamma}{\text{maximize}} \gamma \quad (2)$$

$$\text{subject to: } \sum_{j=1}^p \beta_j^2 = 1 \quad (3)$$

$$\text{and } y_i \left( \beta_0 + \sum_{j=1}^p \beta_j x_{ij} \right) \geq \gamma, \forall i = 1, \dots, N \quad (4)$$

## Constructing the maximal margin classifier (2)

The constraint

$$y_i \left( \beta_0 + \sum_{j=1}^p \beta_j x_{ij} \right) \geq \gamma, \forall i = 1, \dots, N$$

guarantees that each observation is on the correct side of the hyperplane *and* on the correct side of the margin, if margin  $\gamma$  is positive. (This is analogous to Equation 1, but we have added a margin.)

## Constructing the maximal margin classifier (3)

The constraint

$$\text{and } \sum_{j=1}^p \beta_j^2 = 1$$

is not really a constraint: if a separating hyperplane is defined by  $\beta_0 + \sum_{j=1}^p \beta_j x_{ij} = 0$ , then for any  $k \neq 0$ ,  $k \left( \beta_0 + \sum_{j=1}^p \beta_j x_{ij} \right) = 0$  is also a separating hyperplane.

This “constraint” just scales weights so that distance from  $i$ th sample to the hyperplane is given by  $y_i \left( \beta_0 + \sum_{j=1}^p \beta_j x_{ij} \right)$ . This is what make the previous constraint meaningful!

## Constructing the maximal margin classifier (4)

Therefore, the constraints ensure that

- Each observation is on the correct side of the hyperplane, and
- at least  $\gamma$  away from the hyperplane

and  $\gamma$  is maximized.

### Problems with MM classifier (1)




Figure 3: ISLR Fig. 9.4: data may not be separable. Optimization problem has no solution with  $\gamma > 0$ .

## Problems with MM classifier (2)

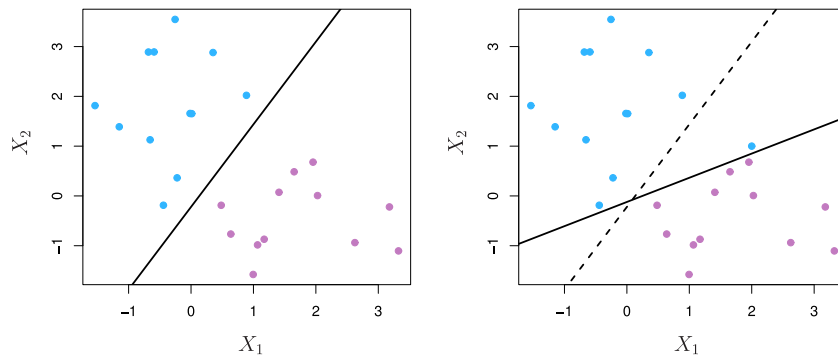


Figure 4: ISLR Fig. 9.5: MM classifier is not robust.

## Support vector classifier

### Basic idea

- Generalization of MM classifier to non-separable case
- Use a hyperplane that *almost* separates the data
- “Soft margin”

### Constructing the support vector classifier

$$\underset{\beta, \epsilon, \gamma}{\text{maximize}} \gamma \quad (5)$$

$$\text{subject to: } \sum_{j=1}^p \beta_j^2 = 1 \quad (6)$$

$$y_i \left( \beta_0 + \sum_{j=1}^p \beta_j x_{ij} \right) \geq \gamma(1 - \epsilon_i), \forall i = 1, \dots, N \quad (7)$$

$$\epsilon_i \geq 0, \sum_{i=1}^N \epsilon_i \leq C \quad (8)$$

$C$  is a non-negative tuning parameter.

### Constructing the support vector classifier (3)

**Slack variable**  $\epsilon_i$  determines where a point lies:

- If  $\epsilon_i = 0$ , point is on the correct side of margin
- If  $\epsilon_i > 0$ , point has *violated* the margin (wrong side of margin)
- If  $\epsilon_i > 1$ , point is on wrong side of hyperplane and is misclassified

## Constructing the support vector classifier (4)

$C$  is the **budget** that determines the number and severity of margin violations we will tolerate.

- $C = 0 \rightarrow$  same as MM classifier
- $C > 0$ , no more than  $C$  observations may be on wrong side of hyperplane
- As  $C$  increases, margin widens; as  $C$  decreases, margin narrows.

## Illustration of effect of $C$




Figure 5: ISLR Fig. 9.7: Margin shrinks as  $C$  decreases.

## Support vector

For a support vector classifier, the only points that affect the classifier are:

- Points that lie on the margin boundary
- Points that violate margin

These are the *support vectors*.

## $C$ controls bias-variance tradeoff

- When  $C$  is large: many support vectors, variance is low, but bias may be high.
- When  $C$  is small: few support vectors, high variance, but low bias.

## Important terminology note

In ISLR and in these notes, meaning of  $C$  is opposite its meaning in Python `sklearn`:

- ISLR and these notes: Large  $C$ , wide margin.
- Python `sklearn`: Large  $C$ , small margin.



### Constrained vs. Lagrange forms

In general, we may see a model expressed in **constrained form**, with tuning parameter  $t \in \mathbb{R}$ :

$$\underset{x \in \mathbb{R}^n}{\text{minimize}} f(x) \text{ subject to } h(x) \leq t$$

and also in **Lagrange form**, with tuning parameter  $\lambda \geq 0$ :

$$\underset{x \in \mathbb{R}^n}{\text{minimize}} f(x) + \lambda h(x)$$

### Loss + penalty expression (1)

Equivalent expression for fitting support vector classifier using *hinge loss*:

$$\underset{\beta}{\text{minimize}} \left( \sum_{i=1}^N \max[0, 1 - y_i f(x_i)] + \lambda \sum_{j=1}^p \beta_j^2 \right)$$

where  $\lambda$  is non-negative tuning parameter similar to  $C$  (large  $\lambda$  means wider margin) and  $f(x_i) = \beta_0 + \sum_{j=1}^p \beta_j x_{ij}$ .

### Loss + penalty representation (2)

With this representation: Zero loss for observations where

$$y_i \left( \beta_0 + \sum_{j=1}^p \beta_j x_{ij} \right) \geq 1$$

and width of margin depends on  $\sum \beta_j^2$ .

### Loss + penalty representation (3)

This is in contrast to previous representation, where: Zero loss for observations where

$$y_i \left( \beta_0 + \sum_{j=1}^p \beta_j x_{ij} \right) \geq \gamma$$

and  $\sum \beta_j^2 = 1$ .

### Compared to logistic regression

- **Hinge loss**: zero for points on correct side of margin.
- **Logistic regression loss**: small for points that are far from decision boundary.

### Hinge loss vs. logistic regression




Figure 6: ISLR 9.12. Hinge loss is zero for points on correct side of margin.

## Maximizing the margin

### Optimization review

Reference: Appendix C.3 of Boyd and Vandenberghe, “Introduction to Applied Linear Algebra”.

### Constrained optimization

Basic formulation of constrained optimization problem:

- **Objective:** Minimize  $f(x)$
- **Constraint(s):** subject to  $g(x) \leq 0$

Find a point  $\hat{x}$  that satisfies  $g(\hat{x}) \leq 0$  and, for any other  $x$  that satisfies  $g(x) \leq 0$ ,  $f(x) \geq f(\hat{x})$ .

### Definition of Lagrangian

Define the Lagrangian as the weighted sum of all constraints:

$$\begin{aligned} L(x, \lambda) &= f(x) + \lambda_1 g_1(x) + \cdots + \lambda_p g_p(x) \\ &= f(x) + g(x)^T \lambda \end{aligned}$$

where  $\lambda$  is the *Lagrange multiplier*.  $g(x)^T \lambda$  “attracts” toward the feasible set, away from the non-feasible set.

### Dual problem (with extra details not shown in class)

Expressed in terms of  $L(x, \lambda)$ , the primal problem is equivalent to

$$\min_x \max_{\lambda \geq 0} L(x, \lambda)$$

The dual problem is

$$\max_{\lambda \geq 0} \min_x L(x, \lambda)$$

### KKT conditions (1)

Under some technical conditions: if  $\hat{x}$  is a local minima, then there is a vector  $\hat{\lambda}$  that satisfies:

$$\frac{\partial L}{\partial x_i}(\hat{x}, \hat{\lambda}) = 0, i = 1 \dots, n$$

$$\frac{\partial L}{\partial \lambda_i}(\hat{x}, \hat{\lambda}) = 0, i = 1 \dots, p$$

(produces as many equations as there are unknowns!)

### KKT conditions (2)

$$g_i(x) \leq 0, \quad i = 1, \dots, p$$

$$\lambda_i \geq 0, \quad i = 1, \dots, p$$

$$\lambda_i g_i(x) = 0, \quad i = 1, \dots, p$$

### Active vs. inactive constraints

At the optimal point, some constraints will be “binding” and some will be “slack” - either:

- $g_i(\hat{x}) < 0$  and  $\hat{\lambda}_i = 0$  (optimum is inside feasible set, constraint is inactive)
- $g_i(\hat{x}) = 0$  and  $\hat{\lambda}_i \geq 0$  (optimum is outside feasible set, constraint is active)

### Active vs. inactive constraints (illustration)




Figure 7: Image via Wikipedia

### Comment on notation

For the following section, we use `sklearn` notation, with opposite meaning of  $C$  -

- in the previous formulation we had a tuning parameter  $\lambda$  that multiplied the penalty term, and increasing this parameter widens the margin
- now  $C$  multiplies the loss term, and increasing this parameter narrows the margin.

### Support vector classifier as constrained optimization (1)

The support vector classifier problem is:

$$\underset{\beta}{\text{minimize}} \left( C \sum_{i=1}^N \epsilon_i + \frac{1}{2} \sum_{j=1}^p \beta_j^2 \right)$$

subject to:

$$y_i(\beta_0 + \sum_{j=1}^p \beta_j x_{ij}) \geq 1 - \epsilon_i \text{ and } \epsilon_i \geq 0, \quad \forall i = 1, \dots, N$$

### Support vector classifier as constrained optimization (extra details)

Construct Lagrange function  $L(\beta, \alpha, \mu)$  where

- $\alpha$  is the vector of Lagrange multipliers for the set of constraints  $y_i(\beta_0 + \sum_{j=1}^p \beta_j x_{ij}) \geq 1 - \epsilon_i$
- $\mu$  is the vector of Lagrange multipliers for the set of constraints  $\epsilon_i \geq 0$

Then the dual problem is:

$$\max_{\alpha, \mu} \min_{\beta} L(\beta, \alpha, \mu)$$

subject to

$$\alpha_i \geq 0, \quad \mu_i \geq 0, \quad \forall i$$

which becomes:

$$\max_{\alpha} \sum_i \alpha_i - \frac{1}{2} \sum_{i,j} \alpha_i \alpha_j y_i y_j \mathbf{x}_i \mathbf{x}_j$$

subject to:

$$\sum_i \alpha_i y_i = 0, \quad C \geq \alpha_i \geq 0, \quad \forall i$$

### Support vector classifier as constrained optimization (2)

Optimal coefficients for  $j = 1, \dots, p$  are:

$$\beta_j = \sum_{i=1}^N \alpha_i y_i x_{ij}$$

where  $\alpha_i$  come from the solution to the dual problem.

### Support vector classifier as constrained optimization (3)

- $\alpha_i > 0$  only when  $x_i$  is a support vector (active constraint).
- Otherwise,  $\alpha_i = 0$  (inactive constraint).

### Support vector classifier as constrained optimization (4)

That leaves  $\beta_0$  - for any  $i$  where  $\alpha_i > 0$ , we can find  $\beta_0$  from

$$\beta_0 = y_i - \sum_{j=1}^p \beta_j x_{ij}$$

### Why solve dual problem?

For high-dimension problems (many features), dual problem can be much faster to solve than primal problem:

- Primal problem: optimize over  $p + 1$  coefficients.
- Dual problem: optimize over  $n$  dual variables, but there are only as many non-zero ones as there are support vectors.

### Correlation interpretation (1)

Given a new sample  $\mathbf{x}$  to classify, compute

$$\hat{z}(\mathbf{x}) = \beta_0 + \sum_{j=1}^p \beta_j x_j = \beta_0 + \sum_{i=1}^N \alpha_i y_i \sum_{j=1}^p x_{ij} x_j$$

Measures inner product (a kind of “correlation”) between new sample and each support vector.

### Correlation interpretation (2)

Classifier output (assuming -1,1 labels):

$$\hat{y}(\mathbf{x}) = \text{sign}(\hat{z}(\mathbf{x}))$$

Predicted label is weighted average of labels for support vectors, with weights proportional to “correlation” of test sample and support vector.

## Support vector machines

### Extension to non-linear decision boundary

- For logistic regression: we used functions of  $\mathbf{x}$  to increase the feature space to classify data that is not linearly separable.
- Could use similar approach here.

### SVM in transformed form (1)

Coefficients:

$$\beta_j = \sum_{i=1}^N \alpha_i y_i \phi(\mathbf{x}_{ij})$$

Classifier discriminant:

$$z = \beta_0 + \sum_{i=1}^N \alpha_i y_i \phi(\mathbf{x}_i) \phi(\mathbf{x})$$

### SVM in transformed form (2)

Classifier output:

$$\hat{y} = \text{sign}(z)$$

**Important:** solution uses inner product of transformed samples, not necessarily transformed samples themselves.

### Kernel trick

$K(\mathbf{x}_i, \mathbf{x}) = \phi(\mathbf{x}_i)\phi(\mathbf{x})$  is a “kernel”.

Classifier discriminant with kernel:

$$z = \beta_0 + \sum_{i=1}^N \alpha_i y_i K(\mathbf{x}_i, \mathbf{x})$$

Can directly compute  $K(\mathbf{x}_i, \mathbf{x})$  **without explicitly computing**  $\phi(\mathbf{x})$ !

(For more details: Mercer’s theorem)

### Kernel trick example

Kernel can be inexpensive to compute, even if basis function itself is expensive. For example, consider:

$$\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix}, \phi(\mathbf{x}) = \begin{bmatrix} x_1^2 \\ x_2^2 \\ \sqrt{2}x_1x_2 \end{bmatrix}$$

### Kernel trick example - direct computation

Direct computation of  $\phi(\mathbf{x}_n)\phi(\mathbf{x}_m)$ : square or multiply 3 components of two vectors (6 operations), then compute inner product in  $\mathbb{R}^3$  (3 multiplications, 1 sum).

$$\begin{aligned} \phi(\mathbf{x}_n)^\top \phi(\mathbf{x}_m) &= \begin{bmatrix} x_{n,1}^2 & x_{n,2}^2 & \sqrt{2}x_{n,1}x_{n,2} \end{bmatrix} \cdot \begin{bmatrix} x_{m,1}^2 \\ x_{m,2}^2 \\ \sqrt{2}x_{m,1}x_{m,2} \end{bmatrix} \\ &= x_{n,1}^2x_{m,1}^2 + x_{n,2}^2x_{m,2}^2 + 2x_{n,1}x_{n,2}x_{m,1}x_{m,2}. \end{aligned}$$

### Kernel trick example - computation using kernel

Using kernel  $K(x_n, x_m) = (x_n^T x_m)^2$ : compute inner product in  $\mathbb{R}^2$  (2 multiplications, 1 sum) and then square of scalar (1 square).

$$\begin{aligned} (\mathbf{x}_m^\top \mathbf{x}_n)^2 &= \left( \begin{bmatrix} x_{n,1} & x_{n,2} \end{bmatrix} \cdot \begin{bmatrix} x_{m,1} \\ x_{m,2} \end{bmatrix} \right)^2 \\ &= (x_{n,1}x_{m,1} + x_{n,2}x_{m,2})^2 \\ &= (x_{n,1}x_{m,1})^2 + (x_{n,2}x_{m,2})^2 + 2(x_{n,1}x_{m,1})(x_{n,2}x_{m,2}) \\ &= \phi(\mathbf{x}_n)^\top \phi(\mathbf{x}_m). \end{aligned}$$

### Kernel intuition

$K(\mathbf{x}_i, \mathbf{x})$  measures “similarity” between training sample  $\mathbf{x}_i$  and new sample  $\mathbf{x}$ .

- Large  $K$ , more similarity

- $K$  close to zero, not much similarity

$z = \beta_0 + \sum_{i=1}^N \alpha_i y_i K(\mathbf{x}_i, \mathbf{x})$  gives higher weight to training samples that are close to new sample.

### Linear kernel




Figure 8: Linear kernel:  $K(x, y) = x^T y$

### Polynomial kernel




Figure 9: Polynomial kernel:  $K(x, y) = (\gamma x^T y + c_0)^d$



## Radial basis function kernel




Figure 10: Radial basis function:  $K(x, y) = \exp(-\gamma\|x - y\|^2)$ . If  $\gamma = \frac{1}{\sigma^2}$ , this is known as the Gaussian kernel with variance  $\sigma^2$ .

## Infinite-dimensional feature space

With kernel method, can operate in infinite-dimensional feature space! Take for example the RBF kernel:

$$K_{\text{RBF}}(\mathbf{x}, \mathbf{y}) = \exp\left(-\gamma\|\mathbf{x} - \mathbf{y}\|^2\right)$$

Let  $\gamma = \frac{1}{2}$  and let  $K_{\text{poly}(r)}$  be the polynomial kernel of degree  $r$ . Then

## Infinite-dimensional feature space (extra steps not shown in class)

$$\begin{aligned} K_{\text{RBF}}(\mathbf{x}, \mathbf{y}) &= \exp\left(-\frac{1}{2}\|\mathbf{x} - \mathbf{y}\|^2\right) \\ &= \exp\left(-\frac{1}{2}\langle \mathbf{x} - \mathbf{y}, \mathbf{x} - \mathbf{y} \rangle\right) \\ &\stackrel{*}{=} \exp\left(-\frac{1}{2}\langle \mathbf{x}, \mathbf{x} - \mathbf{y} \rangle - \langle \mathbf{y}, \mathbf{x} - \mathbf{y} \rangle\right) \\ &\stackrel{*}{=} \exp\left(-\frac{1}{2}\langle \mathbf{x}, \mathbf{x} \rangle - \langle \mathbf{x}, \mathbf{y} \rangle - [\langle \mathbf{y}, \mathbf{x} \rangle - \langle \mathbf{y}, \mathbf{y} \rangle]\right) \\ &= \exp\left(-\frac{1}{2}\langle \mathbf{x}, \mathbf{x} \rangle + \langle \mathbf{y}, \mathbf{y} \rangle - 2\langle \mathbf{x}, \mathbf{y} \rangle\right) \\ &= \exp\left(-\frac{1}{2}\|\mathbf{x}\|^2\right) \exp\left(-\frac{1}{2}\|\mathbf{y}\|^2\right) \exp\left(-2\langle \mathbf{x}, \mathbf{y} \rangle\right) \end{aligned}$$

where the steps marked with a star use the fact that for inner products,  $\langle \mathbf{u} + \mathbf{v}, \mathbf{w} \rangle = \langle \mathbf{u}, \mathbf{w} \rangle + \langle \mathbf{v}, \mathbf{w} \rangle$ .

## Infinite-dimensional feature space (2)

Let  $C$  be a constant

$$C \equiv \exp\left(-\frac{1}{2}\|\mathbf{x}\|^2\right) \exp\left(-\frac{1}{2}\|\mathbf{y}\|^2\right)$$

And note that the Taylor expansion of  $e^{f(x)}$  is:

$$e^{f(x)} = \sum_{r=0}^{\infty} \frac{[f(x)]^r}{r!}$$

Finally, the RBF kernel can be viewed as an infinite sum over polynomial kernels:

$$\begin{aligned} K_{\text{RBF}}(\mathbf{x}, \mathbf{y}) &= C \exp(-2\langle \mathbf{x}, \mathbf{y} \rangle) \\ &= C \sum_{r=0}^{\infty} \frac{\langle \mathbf{x}, \mathbf{y} \rangle^r}{r!} \\ &= C \sum_r \frac{K_{\text{poly}(r)}(\mathbf{x}, \mathbf{y})}{r!} \end{aligned}$$

## Extension to regression

- Similar idea
- Only points outside the margin contribute to final cost

### SVR illustration




Figure 11: Support vector regression.

## Summary: SVM

### Key expression

Discriminant can be computed using an inexpensive kernel function on a small number of support vector points ( $i \in S$  are the subset of training samples that are support vectors):

$$z = \beta_0 + \sum_{i \in S} \alpha_i y_i K(\mathbf{x}_i, \mathbf{x})$$

**Key ideas**

- Defines boundary with greatest separation between classes
- Tuning parameter controls complexity (which direction depends on notation/“meaning” of  $C$ )
- Kernel trick allows efficient extension to higher-dimension space: non-linear decision boundary through transformation of features, but without explicitly computing high-dimensional features.