

# Ensemble methods

Fraida Fund

## Contents

Ensemble methods . . . . .	2
Recap: decision trees . . . . .	2
Ensemble methods - the idea . . . . .	2
Ensemble methods - types (1) . . . . .	2
Ensemble methods - types (2) . . . . .	2
Bagging . . . . .	2
Bagging - background . . . . .	2
Bootstrapping . . . . .	2
Bootstrap aggregation . . . . .	2
Bagging trees . . . . .	2
Variance reduction rule . . . . .	3
Variance reduction rule (general) . . . . .	3
Variance reduction rule (bagged trees) . . . . .	3
Correlated trees . . . . .	4
Random forests . . . . .	4
Bagged trees illustration . . . . .	4
A note on computation . . . . .	4
Boosting . . . . .	5
Boosting - training . . . . .	5
AdaBoost (Adaptive Boosting) . . . . .	5
AdaBoost algorithm . . . . .	5
AdaBoost algorithm (inner loop) . . . . .	5
AdaBoost algorithm (final step) . . . . .	5
Gradient Boosting . . . . .	6
Summary of (selected) ensemble methods . . . . .	6

**Math prerequisites for this lecture:** You should know about:

- Variance of a random variable
- Independence of random variables
- Variance of sum of random variables

## Ensemble methods

### Recap: decision trees

- Let trees grow deep - low bias, high variance
- Don't let trees get deep: low variance, high bias

### Ensemble methods - the idea

Combine multiple **weak learners** - having either high bias or high variance - to create an **ensemble** with better prediction

### Ensemble methods - types (1)

- Combine multiple learners with high **variance** in a way that reduces their variance
- Combine multiple learners with high **bias** in a way that reduces their bias

### Ensemble methods - types (2)

- **Parallel**: build base estimators *independently* and then average their predictions. Combined estimator is usually better than any single base estimator because its *variance* is reduced.
- **Sequential**: (boosting) build base estimators *sequentially* and each one tries to reduce the *bias* of the combined estimator.

## Bagging

### Bagging - background

- Designed for, and most often applied to, decision trees
- Name comes from **bootstrap aggregation**

### Bootstrapping

- Basic idea: Sampling **with replacement**
- Each "bootstrap training set" is *same size* as full training set, and is created by sampling with replacement
- Some samples will appear more than once, some samples not at all

### Bootstrap aggregation

- Create multiple versions  $1, \dots, B$  of training set with bootstrap
- Independently train a model on each bootstrap training set: calculate  $\hat{f}_1(x) \dots, \hat{f}_B(x)$
- Combine output of models by voting (classification) or averaging (regression):

$$\hat{f}_{bag}(x) = \frac{1}{B} \sum_{b=1}^B \hat{f}_b(x)$$

### Bagging trees

- Construct  $B$  trees using  $B$  bootstrapped training sets
- Let the trees grow deep, no pruning
- Each individual tree has low bias, high variance
- Average the prediction of the trees to reduce variance (if independent!)

## Variance reduction rule

$$\text{Var}(\bar{X}) = \text{Var}\left(\frac{1}{n} \sum_{i=1}^n X_i\right) \quad (1)$$

$$= \frac{1}{n^2} \text{Var}\left(\sum_{i=1}^n X_i\right) \quad (2)$$

$$= \frac{1}{n^2} \left( \sum_{i=1}^n \text{Var}(X_i) + 2 \sum_{i < j} \text{Cov}(X_i, X_j) \right) \quad (3)$$

$$= \frac{1}{n^2} \cdot n \text{Var}(X_i) \quad (\text{if } X_i \text{ i.i.d.}) \quad (4)$$

$$= \frac{1}{n} \text{Var}(X_i). \quad (5)$$

where:

1. uses definition of the sample mean:  $\bar{X} = \frac{1}{n} \sum_{i=1}^n X_i$ .
2. uses the scaling rule:  $\text{Var}(aY) = a^2 \text{Var}(Y)$  with  $a = \frac{1}{n}$ .
3. uses variance of sum (+ symmetry of covariance):

$$\text{Var}\left(\sum_i X_i\right) = \sum_i \text{Var}(X_i) + \sum_{\substack{j=1 \\ j \neq i}} \text{Cov}(X_i, X_j) = \sum_i \text{Var}(X_i) + 2 \sum_{i < j} \text{Cov}(X_i, X_j)$$

4. because independence  $\Rightarrow \text{Cov}(X_i, X_j) = 0$  for  $i \neq j$ , and iid means  $\text{Var}(X_i)$  is same for all  $i$ .

## Variance reduction rule (general)

For  $n$  i.i.d. random variables, the variance of their mean decreases with  $n$ :

General RV	
The average	$\bar{X} = \frac{1}{n} \sum_{i=1}^n X_i$
Variance with independence	$\text{Var}(\bar{X}) = \frac{1}{n} \text{Var}(X_i)$

## Variance reduction rule (bagged trees)

For  $B$  **independent** bagged trees, variance decreases with  $B$ :

Bagged Trees	
The average	$\hat{f}_{bag}(x) = \frac{1}{B} \sum_{b=1}^B \hat{f}_b(x)$
Variance with independence	$\text{Var}(\hat{f}_{bag}(x)) = \frac{1}{B} \text{Var}(\hat{f}_b(x))$

...but, they are not really independent!

## Correlated trees

Problem: trees produced by bagging are highly correlated, and:

$$\text{Var}(\bar{X}) = \frac{\sigma^2}{n} [1 + (n-1)\rho]$$

where  $\rho = \text{Corr}(X_i, X_j) = \frac{\text{Cov}(X_i, X_j)}{\sqrt{\text{Var}(X_i) \text{Var}(X_j)}}$

- Imagine there is one feature that is strong predictor, several moderate predictors
- Most/all trees will split on this feature
- Averaging correlated quantities does not reduce variance as much.

## Random forests

Grow many decorrelated trees:

- **Bootstrap**: grow each tree with bootstrap resampled data set.
- **Split-variable randomization**: Force each split to consider *only* a subset of  $m$  of the  $p$  predictors.

Typically  $m = \frac{p}{3}$  but this should be considered a tuning parameter.

## Bagged trees illustration

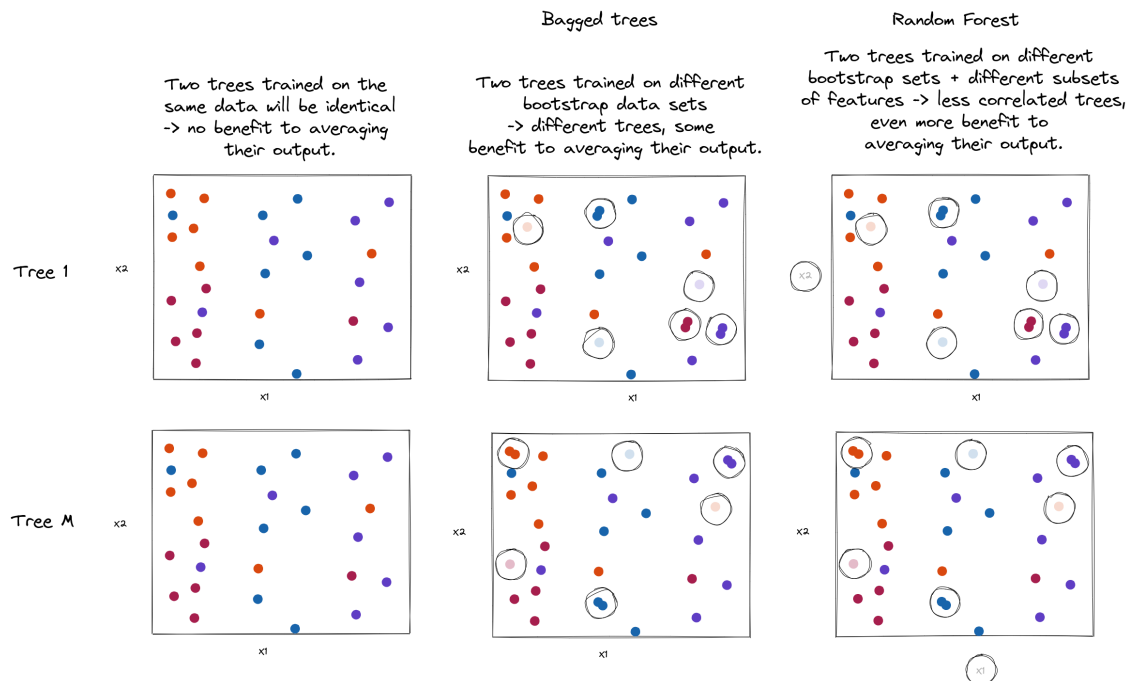


Figure 1: Identical data, bootstrapped data, and bootstrapped data with split variable randomization.

## A note on computation

- Bagged trees and random forests can be fitted in parallel on many cores!
- Each tree is built independently of the others

## Boosting

### Boosting - training

**Iteratively** build a succession of models:

- Train a weak model. Typically a very shallow tree.
- In training set for  $b$ th model, focus on errors made by  $b - 1$ th model.
- Use (weighted) model output
- Reduces bias *and* variance!

### AdaBoost (Adaptive Boosting)

Adjust *weights* so that each successive model focuses on more “difficult” samples.

Consider a binary classification problem, where  $y_i \in \{-1, +1\} \forall i$ .

#### AdaBoost algorithm

1. Let  $w_i = \frac{1}{N}$  for all  $i$  in training set.
2. For  $m = 1, \dots, M$ , repeat:

#### AdaBoost algorithm (inner loop)

- Fit weak learner  $\hat{f}^m$ , on training data with sample weights  $w_i$ .
- Compute weighted error  $err_m$ :

$$err_m = \frac{\sum_{i=1}^N w_i 1(y_i \neq \hat{f}^m(x_i))}{\sum_{i=1}^N w_i}$$

- Compute coefficient  $\alpha_m = \log\left(\frac{1-err_m}{err_m}\right)$
- Update weights:  $w_i \leftarrow w_i e^{\alpha_m 1(y_i \neq \hat{f}^m(x_i))}$  (for misclassified samples, scale weight by  $e^{\alpha_m}$ )

#### AdaBoost algorithm (final step)

3. Output final ensemble model:

$$f_M(x) = \sum_{m=1}^M \alpha_m \hat{f}^m(x), \quad \hat{y}(x) = \text{sign}[f_M(x)]$$

## Gradient Boosting

- General goal of boosting: find the model at each stage that minimizes loss function on ensemble (computationally difficult!)
- AdaBoost interpretation (discovered years later): Gradient descent algorithm that minimizes exponential loss function.
- Gradient boosting: works for any differentiable loss function. At each stage, find the local gradient of loss function, and take steps in direction of steepest descent.

Gradient boosting can be viewed as *functional gradient descent*. We have a differentiable loss function

$$L(y, f(x)),$$

(exponential loss function, in the case of AdaBoost), and an additive model:

$$f_M(x) = \sum_{m=1}^M \alpha_m \hat{f}^m(x),$$

where each weak learner  $\hat{f}^m(x)$  acts like a basis function.

Each iteration performs an update:

$$f_m(x) = f_{m-1}(x) + \alpha_m \hat{f}^m(x).$$

where at each step:

- $\hat{f}^m(x)$  - the weak learner - is chosen to align with the direction of the negative gradient of the loss,
- $\alpha_m$  - the coefficient - determines the step size in that direction.

## Summary of (selected) ensemble methods

- Can use a single estimator that has poor performance
- Combining the output of multiple estimators into a single prediction: better predictive accuracy, less interpretability
- Also more expensive to fit