

Unsupervised learning

Fraida Fund

Contents

Unsupervised learning	3
The basic supervised learning problem	3
The basic unsupervised learning problem	3
Dimensionality reduction with PCA	3
Dimensionality reduction problem	3
Dimensionality reduction with PCA vs feature selection	4
Projections	4
PCA intuition (1)	5
PCA intuition (2)	5
PCA intuition (3)	5
Sample covariance matrix (1)	6
Sample covariance matrix (2)	6
Directional variance	6
Maximizing directional variance (1)	7
Maximizing directional variance (2)	7
Projections onto eigenvectors: uncorrelated features	8
PCA intuition (5)	8
PCA in summary (1)	8
Approximating data	9
Average approximation error	9
Proportion of variance explained	9
Clustering	10
Clustering problem	10
K-means clustering	10
K-means algorithm	10
K-means visualization	10
Dimensionality reduction with deep learning	11
Dimensionality reduction using an autoencoder	11
K-means as an autoencoder (1)	11
K-means as an autoencoder (2)	11
K-means as an autoencoder (3)	11
PCA as an autoencoder (1)	11
PCA as an autoencoder (2)	11
Limits of PCA	11
Neural autoencoder	11
Overcomplete autoencoder	12
Undercomplete autoencoder	12
Sparse autoencoder (1)	13
Sparse autoencoder (2)	13
Autoencoder comparison	13
What are autoencoders good for?	13
Example: reconstruction of faces	14

Example: MNIST visualization	14
Density estimation	15
Types of density estimation	15
GAN: Generative adversarial networks	15
GAN: basic idea (1)	15
GAN: basic idea (2)	15
Discriminator loss function (1)	15
Discriminator loss function (2)	15
Discriminator objective	15
Generator objective (1)	16
Overall objective	16
Problem: gradient of cross-entropy loss	16
Generator objective (2)	16
Training: First, update discriminator	16
Training: Then, update generator	16
Illustration: training discriminator	17
Illustration: training generator	17

Unsupervised learning

The basic supervised learning problem

Given a **sample** with a vector of **features**

$$\mathbf{x} = (x_1, x_2, \dots, x_d)$$

There is some (unknown) relationship between \mathbf{x} and a **target** variable, y , whose value is unknown.

We want to find \hat{y} , our **prediction** for the value of y .

The basic unsupervised learning problem

Given a **sample** with a vector of **features**

$$\mathbf{x} = (x_1, x_2, \dots, x_d)$$

We want to learn something about the underlying *structure* of the data.

No labels!

What are some things we might be able to learn about the structure of the data?

- dimensionality reduction
- clustering
- anomaly detection
- feature learning
- density estimation

Dimensionality reduction with PCA

Why?

- Supervised ML on small feature set
- Visualize data
- Compress data

Dimensionality reduction problem

- Given $N \times p$ data matrix X where each row is a sample x_n
- **Problem:** Map data to $N \times p'$ where $p' \ll p$

Dimensionality reduction with PCA vs feature selection

Previous feature selection:

- Choose subset of existing features
- Many features are somewhat correlated; redundant information

Now: *new* features, so we can get max information with min features.




Figure 1: Instead of using existing features, we project the data onto a new feature space.

Projections

Given vectors z and v , θ is the angle between them. Projection of z onto v is:

$$\hat{z} = \text{Proj}_v(z) = \alpha v, \quad \alpha = \frac{v^T z}{v^T v} = \frac{\|z\|}{\|v\|} \cos \theta$$

$V = \{\alpha v | \alpha \in \mathbb{R}\}$ are the vectors on the line spanned by v , then $\text{Proj}_v(z)$ is the closest vector in V to z : $\hat{z} = \text{argmin}_{w \in V} \|z - w\|^2$.




Figure 2: Projection of z onto v .

PCA intuition (1)




Figure 3: Data with two features, on two axes. Data is centered.

PCA intuition (2)




Figure 4: Construct a new feature by drawing a line $w_1x_1 + w_2x_2$, and projecting data onto that line (red dots are projections). [View animation here.](#)

PCA intuition (3)

Project onto which line?

- Maximize average squared distance from the center to each red dot; **variance of new feature**
- Minimize average squared length of the corresponding red connecting lines; **total reconstruction error**

Can you convince yourself that these two objectives are achieved by the same projection?




Figure 5: Pythagorean decomposition. Keeping reconstruction error minimized (on average) is the same as keeping variance of projection high (on average).

The intuition is that, by Pythagorean decomposition: the variance of the data (a fixed quantity) is equal to the variance of the projected data (which we want to be large) plus the reconstruction error (which we want to be small).

Sample covariance matrix (1)

- sample variance $s_x^2 = \frac{1}{N} \sum_{i=1}^N (x_i - \bar{x})^2$
- sample covariance $s_{xy} = \frac{1}{N} \sum_{i=1}^N (x_i - \bar{x})(y_i - \bar{y})$
- $\text{Cov}(x, y)$ is a $p \times p$ matrix Q with components:

$$Q_{k,l} = \frac{1}{N} \sum_{i=1}^N (x_{ik} - \bar{x}_k)(x_{il} - \bar{x}_l)$$




Figure 6: Illustration of covariance matrix.

Sample covariance matrix (2)

Let \tilde{X} be the data matrix with sample mean removed, row $\tilde{x}_i = x_i - \bar{x}$

Sample covariance matrix is:

$$Q = \frac{1}{N} \tilde{X}^T \tilde{X}$$

(compute covariance matrix by matrix product!)

Now we have these mean-removed rows of data, and we want to project each row onto some vector v , where z is the projection of \tilde{x}_i onto v . And we want to choose v to maximize the variance of z , s_z^2 .

We will call this the *directional variance* - the variance of the projection of the row onto v .

Directional variance

Projection onto v : $z_i = (v^T \tilde{x}_i)v$

- Sample mean: $\bar{z} = v^T \bar{x}$
- Sample variance: $s_z^2 = v^T Q v$

Maximizing directional variance (1)

Given data \tilde{x}_i , what directions of unit vector v ($\|v\| = 1$) maximizes the variance of projection along direction of v ?

$$\max_v v^T Q v \quad \text{s.t.} \|v\| = 1$$

Important note:

- an eigenvector is a special vector that, when you multiply the covariance matrix by the eigenvector, the result is a shorter or longer eigenvector pointing in the *same direction*.
- the eigenvalue is the value by which eigenvector is scaled when multiplied by the covariance matrix.
- a $p \times p$ matrix has p eigenvectors.
- the eigenvectors are orthogonal.

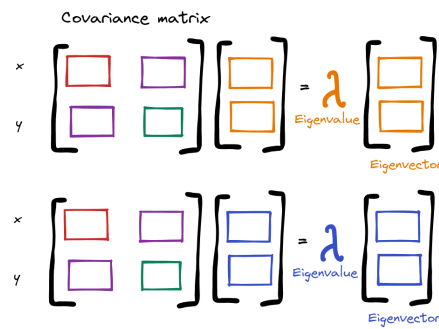


Figure 7: Eigenvectors and eigenvalues.

Maximizing directional variance (2)

Let v_1, \dots, v_p be *eigenvectors* of Q (there are p):

$$Qv_j = \lambda_j v_j$$

- Sort them in descending order: $\lambda_1 \geq \lambda_2 \geq \dots \lambda_p$.
- The largest one is the vector that maximizes directional variance, the next is direction of second most variance, etc.

Theorem: any eigenvector of Q is a local maxima of the optimization problem

$$\max_v v^T Q v \quad \text{s.t.} \|v\| = 1$$

Proof: Define the Lagrangian,

$$L(v, \lambda) = v^T Q v - \lambda [\|v\|^2 - 1]$$

At any local maxima,

$$\frac{\partial L}{\partial v} = 0 \implies Qv - \lambda v = 0$$

Therefore, v is an eigenvector of Q .

Projections onto eigenvectors: uncorrelated features

- Eigenvectors are orthogonal: $v_j^T v_k = 0$ if $j \neq k$
- So the projections of the data onto eigenvectors are uncorrelated

These are called the *principal components*

PCA intuition (5)




Figure 8: In the animation, gray and black lines form a rotating coordinate frame. When variance of projection is maximized, the black line points in direction of first eigenvector of covariance matrix, and grey line points toward second eigenvector. [View animation here.](#)

PCA in summary (1)

Given high-dimensional data,

1. Center data (remove mean)
2. Get covariance matrix
3. Get eigenvectors, eigenvalues
4. Sort by eigenvalue
5. Choose p' eigenvectors with largest eigenvalues
6. Project data onto those eigenvectors

Now you have $N \times p'$ data that maximizes info

Note: in practice, we compute PCA using singular value decomposition (SVD) which is numerically more stable.

Approximating data

Given data $\tilde{x}_i, i = 1, \dots, N$, and PCs v_1, \dots, v_p , we can project + then reconstruct the data:

$$\tilde{x}_i = \sum_{j=1}^p (v_j^T \tilde{x}_i) v_j$$

Consider approximation with *first* $d < p$ coefficients:

$$\hat{x}_i = \sum_{j=1}^d (v_j^T \tilde{x}_i) v_j$$

Average approximation error

For sample i , error is:

$$\tilde{x}_i - \hat{x}_i = \sum_{j=d+1}^p (v_j^T \tilde{x}_i) v_j$$

The projection onto the first principal components carries the most information; the projection onto the last principal components carries the least. So the error due to missing the last PCs is small!

Proportion of variance explained

The *proportion of variance* explained by d PCs is:

$$PoV(d) = \frac{\sum_{j=1}^d \lambda_j}{\sum_{j=1}^p \lambda_j}$$

where the denominator is variance of projected data: $\frac{1}{N} \sum_{i=1}^N ||\tilde{x}_i||^2 = \sum_{j=1}^p \lambda_j$

Clustering

Clustering problem

- Given $N \times d$ data matrix X where each row is a sample x_n
- **Problem:** Group data into K clusters
- More formally: Assign $\sigma_n = \{1, \dots, K\}$ cluster label for each sample
- Samples in same cluster should be close: $\|x_n - x_m\|$ is small when $\sigma_n = \sigma_m$

K-means clustering

We want to minimize

$$J = \sum_{i=1}^K \sum_{n \in C_i} \|x_n - \mu_i\|^2$$

- μ_i is the mean of each cluster
- $\sigma_n \in \{1, \dots, K\}$ is the cluster that x_n belongs to

K-means algorithm

Start with random (?) guesses for each μ_i . Then, iteratively:

- Update cluster membership (nearest neighbor rule): For every n ,

$$\sigma_n = \underset{i}{\operatorname{argmin}} \|x_n - \mu_i\|^2$$

- Update mean of each cluster (centroid rule): for every i , μ_i is average of x_n in C_i
- (Sensitive to initial conditions!)

K-means visualization




Figure 9: Visualization of k-means clustering.

Dimensionality reduction with deep learning

Dimensionality reduction using an autoencoder

An *autoencoder* is a learner that includes:

- Encoder: produces low-dimensional representation of input, $x \rightarrow z$
- Decoder: reconstructs an estimate of input from the low-dimensional representation, $z \rightarrow \hat{x}$
- z known as *latent variables*, *latent representation*, or *code*

K-means as an autoencoder (1)

- Encoder: map each data point to one of K clusters
- Decoder: “reconstruct” data point as center of its cluster

K-means as an autoencoder (2)

- Let $X \in \mathbb{R}^{n \times d}$ be the data matrix containing n d -dimensional data points.
- Let Z be a $n \times k$ matrix (if k clusters) where each column is all zeros, except for one 1
- Let D be a $k \times d$ matrix of cluster centers.

K-means as an autoencoder (3)

- Encoder performs non-linear mapping, expresses result as one-hot vector in Z .
- Decoder is linear:

$$X \approx \hat{X} = ZD$$

PCA as an autoencoder (1)

- Let $X \in \mathbb{R}^{n \times d}$ be the (mean-removed) data matrix containing n d -dimensional data points.
- Let V be a $d \times k$ matrix of k eigenvectors with highest eigenvalues
- $Z = XV$ is the $n \times k$ matrix of PCA projections
- Then $X \approx \hat{X} = ZV^T$

PCA as an autoencoder (2)

- Encoder: linear projection using k best principal components
- Decoder: also linear projection

Limits of PCA

- PCA learns linear projection
- Neural network with non-linear activation function can learn complex non-linear features
- Use neural network to do something like PCA?

Neural autoencoder

- Neural network with d inputs, d outputs
- Use input as target
- (*Self-supervised*: creates its own labels)
- Train network to learn approximation of identity function

What should the architecture of the network be?

Overcomplete autoencoder




Figure 10: If we train this network to minimize reconstruction loss, it may literally learn the identity function - not useful.

Undercomplete autoencoder




Figure 11: This network is forced to learn a low-dimensional representation.

Sparse autoencoder (1)

- Does a small “bottleneck” force autoencoder to learn useful latent features?
- Even if “bottleneck” is very small, can still memorize data by encoding index
- Instead of limiting number of hidden nodes, add a penalty function on *activations*

Sparse autoencoder (2)

Allow many hidden units, but for a given input, most of them must produce a very small activation.

- Add penalty term to loss function, like regularization, but not on weights!
- Penalty is on average activation value (over all the training samples)

Autoencoder comparison

- **Undercomplete autoencoder:** uses entire network for each sample. Limits capacity to memorize, but also limits capacity to extract complex features.
- **Sparse autoencoder:** different parts of network can “specialize” depending on input. Limits capacity to memorize, but can still extract complex features.

What are autoencoders good for?

- Not typically useful for compression - too data-specific
- Can use to initialize supervised learning model - throw away decoder, fine-tune with classifier
- Can use for dimensionality reduction for data visualization (often in combination with other unsupervised learning methods)
- Can use for data denoising

Example: reconstruction of faces




Figure 12: Reconstruction of faces (top) by 30-D neural autoencoder (middle) and 30-D PCA (bottom). Image via Hinton et al “Reducing the dimensionality of data with neural networks”, Science, 2006.

Example: MNIST visualization

Fig. 3. (A) The two-dimensional codes for 500 digits of each class produced by taking the first two principal components of all 60,000 training images. (B) The two-dimensional codes found by a 784-1000-500-250-2 autoencoder. For an alternative visualization, see (8).

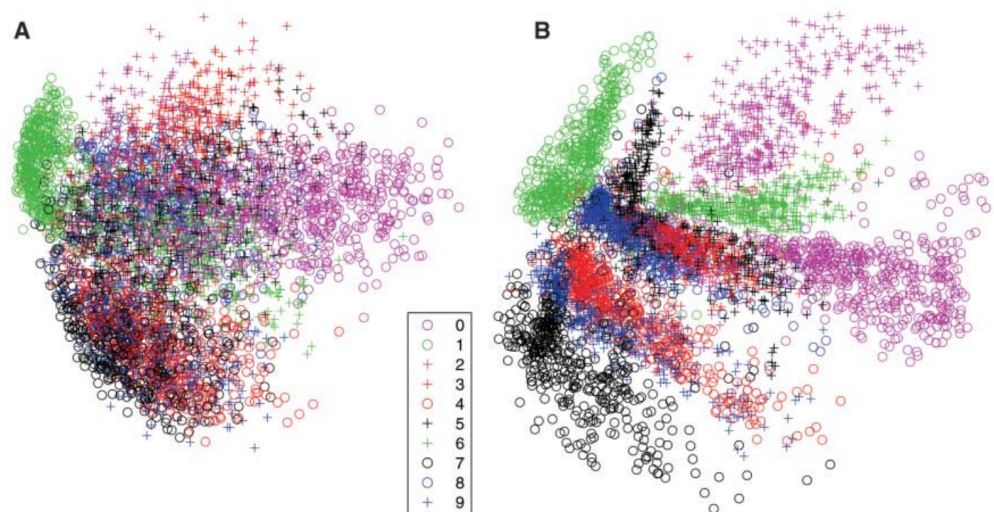


Figure 13: Image via Hinton et al “Reducing the dimensionality of data with neural networks”, Science, 2006.

Density estimation

Types of density estimation

- Explicit: define and solve for density (then sample from it if you want)
- Implicit: sample from density without defining it

GAN: Generative adversarial networks

- From [Goodfellow et al 2014](#)
- Unsupervised, generative, implicit density estimation: Given training data, generate new samples from same distribution

GAN: basic idea (1)

Two neural networks play a “game”:

Generator:

- takes random noise z drawn from p_z as input,
- generates samples, tries to trick “discriminator” into believing they are real,
- learns parameters θ .

GAN: basic idea (2)

Discriminator:

- takes samples x drawn from p_{data} as input,
- produces classification y (1=real, 0=fake),
- learns parameters ϕ .

Discriminator loss function (1)

Discriminator wants to update its parameters ϕ so

- $D_\phi(x)$ (output for real data) is close to 1
- $D_\phi(G_\theta(z))$ (output for generated data) is close to 0

Discriminator loss function (2)

Binary cross-entropy loss:

$$-\sum_{i=1}^N y_i \log D_\phi(x_i) - \sum_{i=1}^N (1 - y_i) \log(1 - D_\phi(x_i))$$

Left side is for “true” samples and the right side is for “fake” samples...

Discriminator objective

Replace sums with expectations, then discriminator wants to *maximize*

$$\mathbb{E}_{x \sim p_{\text{data}}} [\log D_\phi(x)] + \mathbb{E}_{z \sim p_z} [\log(1 - D_\phi(G_\theta(z)))]$$

Generator objective (1)

Generator wants to update its parameters θ so that:

- $D_\phi(G_\theta(z))$ (output for generated data) is close to 1
- Minimize $\mathbb{E}_{z \sim p_z} [\log(1 - D_\phi(G_\theta(z)))]$

Overall objective

$$\min_{\theta} \max_{\phi} \mathbb{E}_{x \sim p_{\text{data}}} [\log D_\phi(x)] + \mathbb{E}_{z \sim p_z} [\log(1 - D_\phi(G_\theta(z)))]$$

Problem: gradient of cross-entropy loss

- Cross-entropy loss designed to accelerate learning (steep gradient) when classifier is wrong
- Gradient is flat when classifier is correct, when generator needs to improve!

Generator objective (2)

- Instead, generator can do gradient *ascent* on the objective

$$\log(D_\phi(G_\theta(z^{(i)})))$$

- Instead of minimizing likelihood of discriminator being correct, now maximizing likelihood of discriminator being wrong.
- Can still learn even when discriminator is successful at rejecting generator samples

Training: First, update discriminator

1. Get mini-batch of size m from data: $x^{(1)}, \dots, x^{(m)} \sim p_{\text{data}}$
2. Get mini-batch of size m from noise input: $z^{(1)}, \dots, z^{(m)} \sim p_z$
3. Forward pass: get $G_\theta(z^{(i)})$ for each noise input, get $D_\phi(x^{(i)})$ for each real sample, get $D_\phi(G_\theta(z^{(i)}))$ for each fake sample.
4. Backward pass: gradient *ascent* on discriminator parameters ϕ :

$$\frac{1}{m} \sum_{i=1}^m [\log D_\phi(x^{(i)}) + \log(1 - D_\phi(G_\theta(z^{(i)})))]$$

Training: Then, update generator

5. Get mini-batch of size m from noise input: $z^{(1)}, \dots, z^{(m)} \sim p_z$
6. Forward pass: get $G_\theta(z^{(i)})$ for each noise input, get $D_\phi(G_\theta(z^{(i)}))$ for each fake sample.
7. Backward pass: gradient *ascent* on generator parameters θ :

$$\frac{1}{m} \sum_{i=1}^m \log(D_\phi(G_\theta(z^{(i)})))$$

Illustration: training discriminator

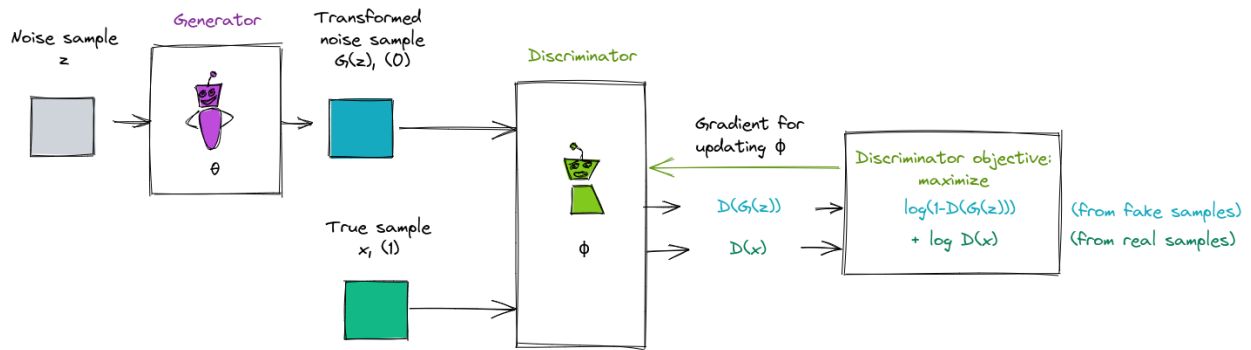


Figure 14: Training the discriminator.

Illustration: training generator




Figure 15: Training the discriminator.