

Intro ML Review

Fraida Fund

Contents

Intro ML review	3
What to know about each model	3
Make decisions about setting it up	3
Compute its output	3
Train for one step	3
What kinds of relationships it can learn	3
What it costs for training/inference	3
What insight we get from it, once trained	3
How to manage its bias/variance	3
What the “pieces” mean	3
What the visualizations mean	3
Not models, but know about...	3
Metrics	3
Evaluation and cross validation	3
Gradient descent	3
Feature selection/weighting	3
Hyperparameter optimization	3
Case studies to review	4
Data leakage case studies	4
Fairness and bias case studies	4
Using learned coefficients case studies	4
Re-thinking “good enough” case studies	4
Moving beyond “out of the box” case studies	4
Code	5
Getting data	5
pandas: read data from file	5
numpy: read data from file	5
pandas to numpy	5
numpy to pandas	5
pandas: stacking columns	5
numpy: stacking columns	5
numpy: reshape 2D to 1D	5
numpy: reshape 1D to 2D	5
numpy: statistics by axis	6
Preprocessing data	6
Ordinal encoding	6
One-hot encoding	6
pandas: converting string to datetime	6
pandas: sort by value of column	6
numpy: sort by value of column	7
pandas: drop missing values	7

numpy: drop missing values	7
pandas: impute missing values	7
numpy: impute missing values	8
numpy: Standardize	8
scikit-learn: Standardize	8
numpy: Min-max scale	8
scikit-learn: Min-max scale	8
scikit-learn: CountVectorizer	8
Create “transformed” features	9
Oversampling/undersampling	9
Slicing and selecting	9
pandas: select columns	9
numpy: select columns	9
pandas: select rows	9
numpy: select rows	9
sklearn: train_test_split	10
sklearn: GroupShuffleSplit	10
sklearn: StratifiedShuffleSplit	10
python: loop with index and value	10
sklearn: KFold with numpy array	10
sklearn: KFold with pandas data frame	10
sklearn: TimeSeriesSplit	11
sklearn: GroupKFold	11
sklearn: StratifiedKFold	11
sklearn: GridSearchCV	11
One-SE rule	11
One-SE rule: Step 1 with lower-is-better metric	11
One-SE rule: Step 1 with higher-is-better metric	12
One-SE rule: Step 2 with models ordered from least complex to most	12
One-SE rule: Step 2 with models ordered from most complex to least	12
Metrics	12
Regression metrics	12
numpy: prediction by mean	12
Classification metrics (label-based)	13
Classification metrics (probability-based)	13
numpy: prediction by mode	13
Model fitting and prediction	14
General pattern	14
Linear models	14
Nearest neighbors	14
Trees and ensembles of trees	14
SVM	14
KMeans clustering	14
PCA	14
Pytorch models	15
Model: Binary classification with prob output	15
Model: Binary classification with logit output	15
Model: Multi-class classification with log prob output	16
Model: Multi-class classification with logit output	16
Model: Regression with one output	16
Model: Regression with multi output	17
Predictions: binary classification with prob output	17
Predictions: binary classification with logit output	17
Predictions: multi-class classification, either output type	17

Predictions: regression, any number of outputs	17
ConvNets: convolution blocks	18
ConvNets: classification head	18
General tips for autograding	18

Intro ML review

(Not comprehensive!)

What to know about each model

Make decisions about setting it up

Compute its output

Train for one step

What kinds of relationships it can learn

What it costs for training/inference

What insight we get from it, once trained

e.g. what its parameters mean

How to manage its bias/variance

What the “pieces” mean

What the visualizations mean

Not models, but know about...

Metrics

Evaluation and cross validation

Gradient descent

Feature selection/weighting

Hyperparameter optimization

Case studies to review

Data leakage case studies

- Beauty in the Classroom
- COVID Case Prediction
- Pre-term Birth Prediction
- Husky vs. Wolf Classification

Fairness and bias case studies

- COMPAS Recidivism Prediction
- Gender Bias in Word Embeddings

Using learned coefficients case studies

- Beauty in the Classroom
- Advertising + PCA Follow-up
- Husky vs. Wolf Classification

Re-thinking “good enough” case studies

- Beauty in the Classroom
- ICU Mortality Prediction

Moving beyond “out of the box” case studies

- Reading a Monkey Brain
- Voter Classification
- UAV-assisted Wireless Localization

Code

Getting data

pandas: **read data from file**

```
df = pd.read_csv("data.csv")
```

```
df = pd.read_table("data.txt", sep="\t")
```

```
df = pd.read_excel("data.xlsx", sheet_name="Sheet1")
```

Common arguments: sep, header, index_col

numpy: **read data from file**

```
X = np.load("X.npy")
```

pandas **to** numpy

```
X = df.to_numpy()
```

```
X = df.values # older way
```

numpy **to** pandas

```
df = pd.DataFrame(X, columns=["a", "b", "c"])
```

pandas: **stacking columns**

```
df = pd.concat([df1, df2], axis=1)
```

(Seen in: L6)

numpy: **stacking columns**

```
X = np.column_stack([X1, X2])
```

numpy: **reshape 2D to 1D**

```
x = X.reshape(-1,)      # flatten to 1D
```

numpy: **reshape 1D to 2D**

```
X = x.reshape(-1, 1)    # shape (n, 1)
```

```
X = x.reshape(1, -1)    # shape (1, n)
```

numpy: statistics by axis

```
m = np.mean(X, axis=0)      # mean of each column
```

```
m = np.mean(X, axis=1)      # mean of each row
```

Common statistics: std, min, max, quantile...

Preprocessing data

Ordinal encoding

```
map_dict = {'18-29': 1, '30-44': 2, '45-64': 3, '65+': 4}
df_enc_ord = pd.DataFrame(
    {'age': df['age'].map(map_dict)},
    index=df.index
)
```

(Seen in: L6)

One-hot encoding

```
# encode one column
df2 = pd.get_dummies(df["col"], dtype=np.int32)
```

(Seen in: L6)

pandas: converting string to datetime

```
df["date"] = pd.to_datetime(df["date"])
```

```
df["date"] = pd.to_datetime(df["date"], format="%Y-%m-%d")
```

(Seen in: Week 1 Exploratory Data Analysis Colab lesson + a few times in HW)

pandas: sort by value of column

```
df2 = df.sort_values(by="col")
```

```
df2 = df.sort_values(by=["col1", "col2"])
```

Common arguments: ascending

(Seen in: Week 1 Exploratory Data Analysis Colab lesson, Week 4 Model Selection Colab lesson, H4 K-fold CV with Fourier basis expansion)

numpy: sort by value of column

```
idx = np.argsort(X[:, 0])
X2 = X[idx]
```

```
# sort by col0, then col1
idx = np.lexsort((X[:, 1], X[:, 0]))
X2 = X[idx]
```

(Seen in: L6 when breaking ties)

pandas: drop missing values

```
# drop rows with ANY missing values
df2 = df.dropna()
```

```
# drop rows missing specific column by name
df2 = df.dropna(subset=["col1"])
```

```
# drop rows missing in either column
df2 = df.dropna(subset=["col1", "col2"])
```

(Seen in: L6 when computing feature weights, H4 TimeSeriesSplit question)

numpy: drop missing values

```
# drop rows with ANY missing values
mask = np.isnan(X).sum(axis=1) == 0
X2 = X[mask]
```

```
# drop rows missing in either column by index
cols = [0, 2]
mask = np.isnan(X[:, cols]).sum(axis=1) == 0
X2 = X[mask]
```

pandas: impute missing values

```
# forward fill, data must be sorted
df["col"] = df["col"].fillna(method="ffill")
```

```
# using a statistic, only use stats of training set
s = df_tr["col"].mean() # use training data only
df_tr["col"] = df_tr["col"].fillna(s)
df_ts["col"] = df_ts["col"].fillna(s)
```

```
# using a constant
df["col"] = df["col"].fillna(0)
```

(Seen in: Week 1 Exploratory Data Analysis Colab lesson, H5 ICU mortality prediction)

numpy: impute missing values

```
# applies to all columns
# using a statistic, only use stats of training set
s = np.nanmean(Xtr, axis=0) # use training data only
Xtr2 = np.where(np.isnan(Xtr), s, Xtr)
Xts2 = np.where(np.isnan(Xts), s, Xts)

# applies to all columns
# using a constant
X2 = np.where(np.isnan(X), 0, X)
```

numpy: Standardize

```
x_mean = np.mean(X[idx_tr], axis=0)
x_std = np.std(X[idx_tr], axis=0)
Xtr_std = (X[idx_tr] - x_mean) / x_std
Xts_std = (X[idx_ts] - x_mean) / x_std
```

scikit-learn: Standardize

```
scaler = StandardScaler()
Xtr_std = scaler.fit_transform(X[idx_tr])
Xts_std = scaler.transform(X[idx_ts])
```

(Seen in: Week 4 Regularization Colab lesson)

numpy: Min-max scale

```
x_min = np.min(X[idx_tr], axis=0)
x_max = np.max(X[idx_tr], axis=0)
Xtr_mm = (X[idx_tr] - x_min) / (x_max - x_min)
Xts_mm = (X[idx_ts] - x_min) / (x_max - x_min)
```

scikit-learn: Min-max scale

```
scaler = MinMaxScaler()
Xtr_mm = scaler.fit_transform(X[idx_tr])
Xts_mm = scaler.transform(X[idx_ts])
```

(Seen in: L6)

scikit-learn: CountVectorizer

```
vect = CountVectorizer()
Xtr = vect.fit_transform(text_tr)      # fit on training text
Xts = vect.transform(text_ts)         # transform test text
```

Common arguments: stop_words

(Seen in: H5)

Create “transformed” features

```
df = df.assign(interaction = df["col1"] * df["col2"])
```

```
# interaction of col0 and col1
interaction = X[:, 0] * X[:, 1]
X_new = np.column_stack([X, interaction])
```

(Seen in: L2)

Oversampling/undersampling

```
ovr = ADASYN(n_neighbors = 5, random_state = random_state)
X_ovr, y_ovr = ovr.fit_resample(X, y)
```

(Seen in: H8)

Slicing and selecting

pandas: select columns

```
col = df["col"]          # one column (Series)
cols = df[["col1", "col2"]]  # list of columns (DataFrame)
```

numpy: select columns

```
x = X[:, 3]          # one column by index
x = X[:, 2:5]         # contiguous columns (2,3,4)
x = X[:, [0, 3, 7]]   # list of columns
```

pandas: select rows

```
# by integer position - use .iloc
rows = df.iloc[5]
```

```
# by condition
rows = df[df["cola"] > 0]
rows = df[(df["cola"] > 0) & (df["colb"] < 5)]
```

numpy: select rows

```
r = X[5]          # one row
r = X[10:20]       # slice of rows
r = X[[0, 3, 7]]   # list of row indices
```

```
# by condition on column 0
mask = X[:, 0] > 0
X2 = X[mask]
```

```
# by condition using argwhere
idx = np.argwhere(X[:, 0] > 0).reshape(-1)
X2 = X[idx]
```

sklearn: train_test_split

```
# split X and y together
Xtr, Xts, ytr, yts = train_test_split(
    X, y, test_size=0.2, random_state=0, shuffle=True
)
```

```
# split X only
Xtr, Xts = train_test_split(X, test_size=100, shuffle=False)
```

Common arguments: test_size (ratio or number), shuffle, random_state

sklearn: GroupShuffleSplit

```
gss = GroupShuffleSplit(test_size=0.2, random_state=0)
train_idx, test_idx = next(gss.split(X, y, groups))
```

sklearn: StratifiedShuffleSplit

```
sss = StratifiedShuffleSplit(test_size=0.2, random_state=0)
train_idx, test_idx = next(sss.split(X, y))
```

python: loop with index and value

```
for i, x in enumerate(X):
    # do stuff with i = index, x = row (X is 2D)
```

sklearn: KFold with numpy array

```
kf = KFold(n_splits=5, shuffle=True, random_state=0)

for i, (idx_tr, idx_ts) in enumerate(kf.split(X)):
    Xtr, Xts = X[idx_tr], X[idx_ts]
```

sklearn: KFold with pandas data frame

```
kf = KFold(n_splits=5, shuffle=True, random_state=0)

for i, (idx_tr, idx_ts) in enumerate(kf.split(X)):
    Xtr, Xts = X.iloc[idx_tr], X.iloc[idx_ts]
```

Note the use of iloc! (Applies to all the “variants,” too.)

sklearn: TimeSeriesSplit

```
tscv = TimeSeriesSplit(n_splits=5)

for i, (idx_tr, idx_ts) in enumerate(tscv.split(X)):
    Xtr, Xts = X[idx_tr], X[idx_ts]
```

sklearn: GroupKFold

```
gkf = GroupKFold(n_splits=5)

for i, (idx_tr, idx_ts) in enumerate(gkf.split(X, groups)):
    Xtr, Xts = X[idx_tr], X[idx_ts]
```

Comment: groups by groups

sklearn: StratifiedKFold

```
skf = StratifiedKFold(n_splits=5, shuffle=True, random_state=0)

for i, (idx_tr, idx_ts) in enumerate(skf.split(X, y)):
    Xtr, Xts = X[idx_tr], X[idx_ts]
```

Comment: stratifies by y

sklearn: GridSearchCV

```
param_grid = [
    {'C': [1e-1, 1e1, 1e3], 'gamma': [1e2, 1, 1e-1, 1e-3]},
]
grid = GridSearchCV(model, param_grid, cv=5)
grid.fit(Xtr, ytr)
best = grid.best_params_
```

(Seen in: H8)

One-SE rule

- **Step 1:** Find models within one SE of “best” model
- **Step 2:** Pick the simplest model from the list of candidates

One-SE rule: Step 1 with lower-is-better metric

```
# score_val is a lower-is-better metric
# so threshold is: SMALLEST mean score PLUS its std/(nfold-1)
idx_min = np.argmin(score_val.mean(axis=1))
target = score_val[idx_min,:].mean() + score_val[idx_min,:].std()/np.sqrt(nfold-1)
# candidate models have mean score BELOW that threshold
idx_one_se = np.where(score_val.mean(axis=1) <= target)
```

One-SE rule: Step 1 with higher-is-better metric

```
# score_val is a higher-is-better metric
# so threshold is: LARGEST mean score MINUS its std/(nfold-1)
idx_min = np.argmax(score_val.mean(axis=1))
target = score_val[idx_min,:].mean() - score_val[idx_min,:].std()/np.sqrt(nfold-1)
# candidate models have mean score ABOVE that threshold
idx_one_se = np.where(score_val.mean(axis=1) >= target)
```

One-SE rule: Step 2 with models ordered from least complex to most

```
# get LOWEST indexed model (least complex)
m_one_se = np.min(model_list[idx_one_se])
```

One-SE rule: Step 2 with models ordered from most complex to least

```
# get HIGHEST indexed model (least complex)
m_one_se = np.max(model_list[idx_one_se])
```

Metrics

Regression metrics

```
mse = mean_squared_error(y, y_pred)
```

```
mae = mean_absolute_error(y, y_pred)
```

```
# r2_score(y_true, y_pred) - order matters: true first, pred second
r2 = r2_score(y, y_pred)
```

Comment: y and y_pred must have same shape...

numpy: prediction by mean

```
# ones * mean, works by broadcasting
m = np.mean(ytr)
pred_mean = np.ones(y.shape) * m
```

```
# fill array with mean
m = np.mean(ytr)
pred_mean = np.full_like(y, m)
```

```
# repeat mean
m = np.mean(ytr)
pred_mean = np.repeat(m, len(y))
```

Comment: use training data only to compute mean! then can fill all of y or just ytr or yts.

Classification metrics (label-based)

```
acc    = accuracy_score(y, y_pred)
bacc   = balanced_accuracy_score(y, y_pred)
f1     = f1_score(y, y_pred)
prec   = precision_score(y, y_pred)
rec    = recall_score(y, y_pred)
```

Comment: y and y_pred must have same shape...

Classification metrics (probability-based)

```
# y_proba = probability for class 1
y_proba = model.predict_proba(X)[:, 1]
auc    = roc_auc_score(y, y_proba)
```

numpy: prediction by mode

Same idea as “prediction by mean”, but to get mode...

```
vals, counts = np.unique(ytr, return_counts=True)
# index of most common value in counts array
m_idx = np.argmax(counts)
# actual label
m = vals[m_idx]
```

```
from scipy.stats import mode
m = mode(ytr, keepdims=True).mode[0]
```

Comment: use training data only to compute mode!

(Seen in: H5 ICU Mortality Prediction, and others)

Model fitting and prediction

General pattern

```
# choose model + pass arguments
model = Model()
# train, Xtr MUST be 2D (n_samples, n_features)
model.fit(Xtr, ytr)
# predict
y_pred = model.predict(Xts)
```

(Seen in: Week 2 Colab lesson)

Linear models

```
model = LinearRegression()
model = Ridge(alpha=1.0)
model = Lasso(alpha=0.1)
model = LogisticRegression()
```

Nearest neighbors

```
model = KNeighborsRegressor(n_neighbors=5)
model = KNeighborsClassifier(n_neighbors=5)
```

Trees and ensembles of trees

```
model = DecisionTreeRegressor(max_depth=5)
model = DecisionTreeClassifier(max_depth=5)

model = RandomForestRegressor(n_estimators=200)
model = RandomForestClassifier(n_estimators=200)

model = AdaBoostClassifier(n_estimators=200, learning_rate=0.5)
```

SVM

```
model = SVC(kernel="rbf", C=1.0, gamma=0.1)
```

KMeans clustering

```
model = KMeans(n_clusters=3, random_state=0)
Xtr_cls = model.fit_predict(Xtr)          # unsupervised: no y
Xts_cls = model.predict(Xts) # cluster assignments on new data
```

PCA

```
pca = PCA(n_components=2)
Xtr_pca = pca.fit_transform(Xtr)      # fit PCA on training data
Xts_pca = pca.transform(Xts)        # apply to test data
```

Pytorch models

Model: Binary classification with prob output

```
nout = 1 # !!!  
  
class SimpleNNBinaryClassification(nn.Module):  
  
    def __init__(self, nin, nh, nout):  
        super(SimpleNNBinaryClassification, self).__init__()  
        self.hidden = nn.Linear(nin, nh)  
        self.output = nn.Linear(nh, nout)  
  
    def forward(self, x):  
        x = torch.sigmoid(self.hidden(x))  
        x = torch.sigmoid(self.output(x)) # !!!  
        return x  
  
model = SimpleNNBinaryClassification(nin, nh, nout)  
loss_fn = nn.BCELoss() # !!!  
optimizer = optim.Adam(model.parameters(), lr=0.01)
```

Model: Binary classification with logit output

```
nout = 1 # !!!  
  
class SimpleNNBinaryClassificationLogits(nn.Module):  
  
    def __init__(self, nin, nh, nout):  
        super(SimpleNNBinaryClassificationLogits, self).__init__()  
        self.hidden = nn.Linear(nin, nh)  
        self.output = nn.Linear(nh, nout)  
  
    def forward(self, x):  
        x = torch.sigmoid(self.hidden(x))  
        x = self.output(x) # !!!  
        return x  
  
model = SimpleNNBinaryClassificationLogits(nin, nh, nout)  
loss_fn = nn.BCEWithLogitsLoss() # !!!  
optimizer = optim.Adam(model.parameters(), lr=0.01)
```

Model: Multi-class classification with log prob output

```
nout = 3 # !!!  
  
class SimpleNNMultiClassification(nn.Module):  
  
    def __init__(self, nin, nh, nout):  
        super(SimpleNNMultiClassification, self).__init__()  
        self.hidden = nn.Linear(nin, nh)  
        self.output = nn.Linear(nh, nout)  
  
    def forward(self, x):  
        x = torch.sigmoid(self.hidden(x))  
        x = torch.log_softmax(self.output(x), dim=1) # !!!  
        return x  
  
model = SimpleNNMultiClassification(nin, nh, nout)  
loss_fn = nn.NLLLoss() # !!!  
optimizer = optim.Adam(model.parameters(), lr=0.01)
```

Model: Multi-class classification with logit output

```
nout = 3 # !!!  
  
class SimpleNNMultiClassificationLogits(nn.Module):  
  
    def __init__(self, nin, nh, nout):  
        super(SimpleNNMultiClassificationLogits, self).__init__()  
        self.hidden = nn.Linear(nin, nh)  
        self.output = nn.Linear(nh, nout)  
  
    def forward(self, x):  
        x = torch.sigmoid(self.hidden(x))  
        x = self.output(x) # !!!  
        return x  
  
model = SimpleNNMultiClassificationLogits(nin, nh, nout)  
loss_fn = nn.CrossEntropyLoss() # !!!  
optimizer = optim.Adam(model.parameters(), lr=0.01)
```

Model: Regression with one output

```
nout = 1 # !!!  
  
class SimpleNNRegression(nn.Module):  
  
    def __init__(self, nin, nh, nout):  
        super(SimpleNNRegression, self).__init__()  
        self.hidden = nn.Linear(nin, nh)  
        self.output = nn.Linear(nh, nout)  
  
    def forward(self, x):  
        x = torch.sigmoid(self.hidden(x))
```

```

        x = self.output(x) # !!!
        return x

model = SimpleNNRegression(nin, nh, nout)
loss_fn = nn.MSELoss() # !!!
optimizer = optim.Adam(model.parameters(), lr=0.01)

```

Model: Regression with multi output

```

nout = 2 # !!!

class SimpleNNMultiRegression(nn.Module):

    def __init__(self, nin, nh, nout):
        super(SimpleNNMultiRegression, self).__init__()
        self.hidden = nn.Linear(nin, nh)
        self.output = nn.Linear(nh, nout)

    def forward(self, x):
        x = torch.sigmoid(self.hidden(x))
        x = self.output(x) # !!!
        return x

model = SimpleNNMultiRegression(nin, nh, nout)
loss_fn = nn.MSELoss() # !!!
optimizer = optim.Adam(model.parameters(), lr=0.01)

```

Predictions: binary classification with prob output

```

with torch.no_grad():
    outputs = model(X)
    preds = (outputs > 0.5).int()

```

Predictions: binary classification with logit output

```

with torch.no_grad():
    logits = model(X)
    outputs = torch.sigmoid(logits)
    preds = (outputs > 0.5).int()

```

Predictions: multi-class classification, either output type

```

with torch.no_grad():
    outputs = model(X)
    preds = torch.argmax(outputs, dim=1)

```

Predictions: regression, any number of outputs

```
with torch.no_grad():
    outputs = model(X)
    preds = outputs
```

ConvNets: convolution blocks

```
# Before: (batch_size, n1, H1, W1)
nn.Conv2d(n1, n2, kernel_size=3, padding=0),
# After: (batch_size, n2, H2, W2)
nn.ReLU(),           # no change in shape
nn.BatchNorm2d(n2), # no change in shape
nn.MaxPool2d(kernel_size=2),
# After: (batch_size, n2, H3, W3)
```

ConvNets: classification head

```
self.classifier = nn.Sequential(
    # Before: (batch_size, n, h, w)
    nn.Flatten(),
    # After: (batch_size, n x h x w)
    nn.Linear(n_out_conv, n_out),
    # After: (batch_size, n_out)
)
```

General tips for autograding

- run your code beginning to end before submitting to grader
- don't re-use out-of-loop variable names inside loops
- keep an eye on #grade tags